

Practical R for MBM

Chuck Frost

7/10/2020

Practical R for MBM field data manipulation and analysis

The purpose of this ongoing short course is to familiarize MBM personnel with the concepts of tidy data, basic data manipulation, and basic programming techniques as implemented in R. It is expected that participants have a working understanding of R and RStudio. We will not go into depth on other types of programming, specifically statistical programming and computer programming. The topics covered here will foster appropriate treatment of MBM data throughout the data life cycle and lead to cleaner and more useful data in the future.

Topics

1. R, RStudio, packages, help, intro to functions
2. Loading and saving, data types and structures
3. Tidy Data
4. Visualization
5. Creating functions, conditional statements, looping
6. Working with dates

1.0.1 What is R?

An open-source programming environment that serves as:

1. A giant calculator
2. An interface for data analysis and visualization
3. An interface for data manipulation and file handling
4. An interface for a simple and efficient programming language (S, that became R)

R isn't the "method." For example, "We used R to estimate population size." yuck

1.0.2 What is RStudio?

An integrated development environment (IDE) for R.

Provides tools, menus, help, and easy access to the best parts of R.

1.1.1 Packages

Base R automatically provides you with many common functions.

```
mean(c(1,2,3))
```

```
## [1] 2
```

```
sum(c(1,2,3))
```

```
## [1] 6
```

```
var(c(1,2,3))
```

```
## [1] 1
```

Other useful functions exist thanks to R users. Collections of related functions get wrapped into packages.

If you're looking for something NOT in base R, you need to install the related package.

The CRAN provides R community tested and approved packages. Check the packages tab in RStudio for some you already have installed but (probably) not loaded. Just click them to load. If you know what package you need, there are 2 easy ways to install it. The first is by clicking the install button in the packages tab. Search the CRAN for the package you want and install it. The second is:

```
install.packages("PackageName")
```

Don't forget to load the package after installation, either by checking the box under packages, or by calling:

```
library(PackageName)
```

Note that in `install.packages()` you are searching a string in quotes, "PackageName" while in `library()` you are calling a package object without quotes, PackageName.

1.1.2 Installing from GitHub

If a package isn't uploaded to CRAN, it can still be made available through GitHub. Packages installed from GitHub require an intermediate step:

```
install.packages("devtools")
```

```
library(devtools)
```

```
install_github("USFWS/AKaerial", ref = "development")
```

```
library(AKaerial)
```

This code does 4 things:

1. Install devtools package
2. Load devtools package

3. Call `install_github` (from `devtools`), look for the GitHub account “USFWS” and repository “AKaerial” and load the package in the “development” branch
4. Load AKaerial package

Installing and loading AKaerial may take a while since it includes dependencies. Dependencies are packages that are specified within another package that must be included for that package to operate appropriately. AKaerial “depends” on several other packages and will check if you have these installed.

1.1.3 Functions

Once you have a package loaded, you can access functions written and included in that package. Typing the name of a function will spit out the code that makes up the function. This can be messy for a long or complicated function. Typing `?FunctionName` will access the help file for a given function.

```
AdjustCounts
```

```
?AdjustCounts
```

This can be extremely helpful (if the help file is!) when troubleshooting why your code isn’t working as you want it to. Note that the top of the help file also tells you `FunctionName {PackageName}` in case you can’t figure out where your function is coming from. This could be the case if you have dozens of packages loaded or have borrowed someone else’s code to help run your analysis. Don’t underestimate the power of internet searches to find if something already exists to help you do your thing!

We will talk about functions in more depth later, but for now, the general idea is that functions take arguments, run processes, and (usually) return products saved as R objects.

```
numbers = c(1,2,3)
```

```
avg = mean(numbers)
```

```
avg
```

```
## [1] 2
```

In this example, `numbers` is initiated as a vector of 3 integers: 1,2,3. We then run the function `mean` with `numbers` as its argument, and save the result as the object `avg`.

It can be helpful to think of functions as recipes and arguments as ingredients.

```
Mix = function(what.to.mix){  
  dough = sum(what.to.mix)  
  return(dough)  
}
```

```
Bake = function(what.to.bake, time.to.bake){  
  
  cookies = what.to.bake * time.to.bake  
  
  return(cookies)  
}
```

```
my.ingredients = c("sugar", "eggs", "butter", "flour")

my.dough = Mix(my.ingredients)

my.cookies = Bake(my.dough, 10)
```

2.1.0 Loading Files

There are many ways to read data into R. We will go over the most common way data is stored and used in R; as a data frame. A data frame is a 2-dimensional array (table) where each column represents a variable (category) and each row represents a unique observation. Let's start with creating your own data frame. This might be useful if you have a small data set.

```
small.data = data.frame(
  name = c("Laura", "Zak", "Chuck"),
  size = c("small", "medium", "large"),
  tenure = c(6, "drifter", 9)
)
```

We can now check the structure of the data frame. This is a great first QA/QC check! If we know that tenure should be numeric, the structure should confirm it.

```
str(small.data)
```

```
## 'data.frame':  3 obs. of  3 variables:
## $ name  : Factor w/ 3 levels "Chuck","Laura",...: 2 3 1
## $ size  : Factor w/ 3 levels "large","medium",...: 3 2 1
## $ tenure: Factor w/ 3 levels "6","9","drifter": 1 3 2
```

Notice how all 3 variables are shown as Factor, which is incorrect. Factors, in general are tough to work with. If we try to change the value here (use \$ to grab a column and brackets to select a row [], or use only brackets to select [row, column]), it throws an error.

```
small.data$tenure[2] = 3
```

```
## Warning in '[<-factor'('*tmp*', 2, value = structure(c(1L, NA, 2L), .Label
## = c("6", : invalid factor level, NA generated
```

```
small.data[2,3] = 3
```

```
## Warning in '[<-factor'('*tmp*', iseq, value = 3): invalid factor level, NA
## generated
```

```
levels(small.data$tenure)
```

```
## [1] "6"      "9"      "drifter"
```

Since R couldn't determine the structure of our tenure column, it defaulted to calling it a Factor, which can take any value. The unique set of values gets set as the levels of that factor. If you try to define a value outside of that set, you will get an error that results in NA. This is (almost) never what you want. So we can fix this in 2 ways:

First, we can use the "as." functions in base R to force R to evaluate something as a different data type. In this case, we can't just let R evaluate the factor as numeric since it will then take a numeric representation of that particular level of a factor in the set of all levels of that factor. What a mess. We have to first convert the factor to a character string using `as.character`, then to a numeric using `as.numeric`. This is called wrapping functions, when you pass as an argument to a new function the output of another function without saving the intermediate result. In the outer-most function call to `as.numeric`, anything that R can't represent as numeric will get NA. This is already messy and hard to trace changes to your data!

```
small.data$tenure=as.numeric(as.character(small.data$tenure))

str(small.data)
```

```
## 'data.frame':   3 obs. of  3 variables:
## $ name   : Factor w/ 3 levels "Chuck","Laura",...: 2 3 1
## $ size   : Factor w/ 3 levels "large","medium",...: 3 2 1
## $ tenure: num  6 NA 9
```

Or, since we are scripting our analysis, we can change it in our original code chunk and just re-run. Score 1 for a nice scripted workflow...but imagine doing this through thousands or hundreds of thousands of rows of data.

```
small.data = data.frame(
  name = c("Laura", "Zak", "Chuck"),
  size = c("small", "medium", "large"),
  tenure = c(6,3,9)
)
```

While we are at it, let's just purge the factors from our data entirely. It is easy to make something a factor later, if we ever see the need. To do this, we just pass the argument `stringsAsFactors = FALSE` to our data frame function. Note that we have to add a comma after the declaration of our final column tenure.

```
small.data = data.frame(
  name = c("Laura", "Zak", "Chuck"),
  size = c("small", "medium", "large"),
  tenure = c(6,3,9),
  stringsAsFactors = FALSE
)

str(small.data)
```

```
## 'data.frame':   3 obs. of  3 variables:
## $ name   : chr  "Laura" "Zak" "Chuck"
## $ size   : chr  "small" "medium" "large"
## $ tenure: num  6 3 9
```

Now we are clean...maybe. Our structure says our name and size columns are chr (character) and our tenure correctly shows (num) numeric. Now we want add a new column (greatness) using the values in an existing column. We can both create and define the values of the new column in one step.

```
small.data$greatness = small.data$tenure^3

str(small.data)
```

```
## 'data.frame': 3 obs. of 4 variables:
## $ name : chr "Laura" "Zak" "Chuck"
## $ size : chr "small" "medium" "large"
## $ tenure : num 6 3 9
## $ greatness: num 216 27 729
```

```
small.data
```

```
##   name   size tenure greatness
## 1 Laura small     6      216
## 2  Zak medium     3       27
## 3 Chuck large     9      729
```

Now we have our 4 columns and the structure even gives us what it can display of the values. Now what if we want to add another row of data? The function `rbind` (row bind) makes it easy. As long as our new row is in the same order and type as the columns of our data frame, we can simply call:

```
small.data = rbind(small.data, c("Hannah", "small", 1, 1))

str(small.data)
```

```
## 'data.frame': 4 obs. of 4 variables:
## $ name : chr "Laura" "Zak" "Chuck" "Hannah"
## $ size : chr "small" "medium" "large" "small"
## $ tenure : chr "6" "3" "9" "1"
## $ greatness: chr "216" "27" "729" "1"
```

```
small.data
```

```
##   name   size tenure greatness
## 1 Laura small     6      216
## 2  Zak medium     3       27
## 3 Chuck large     9      729
## 4 Hannah small     1        1
```

As you can imagine, the same thing can be done for columns using `cbind`. Let's add another one:

```
small.data = cbind(small.data, interests=c("Bedazzling", "K-Pop", "Couponing", "Oranges"))

small.data
```

```
##   name   size tenure greatness interests
## 1 Laura small     6      216 Bedazzling
## 2  Zak medium     3       27      K-Pop
## 3 Chuck large     9      729 Couponing
## 4 Hannah small     1        1    Oranges
```

2.2.0 Reading in a data file

If you have a data file that exists (hopefully) as .csv or .txt, base R has a function that will read your data into a data frame for you. We will demo this with a .csv of fake aerial survey data. Remember that R works best (and sometimes only!) with tidy data. Your file should at least be rectangular (no loose columns, no long or dangling rows) and preferably will have column headings, though those can also be defined later. The function `read.table` will be the most generic and robust method, but there is also `read.csv` that saves a little time since you are predefining the formatting.

Plain text format is the simplest way to store text. One step up are delimited files. The most common types of delimited files are tab- and comma-delimited. Comma-delimited files are commonly saved as .csv (comma separated values). R does not care how your files are delimited, you just have to know the delimiting character(s) and pass them as arguments in the function call.

```
my.data = read.table("SomeData.txt", header = TRUE, sep = " ")
```

This tells R to find the file `SomeData.txt` (it will default to your working directory, but you can path it anywhere). It will open it, read the header row into the data frame `my.data`, then scan the file for spaces (`sep = " "`) and add entries between spaces to the values in each row. What could possibly go wrong with space-delimiting? Or even comma-delimiting?

To illustrate further concepts below we will be using an artificial data set I modified for this demo. It is located on the GitHub repository for this short course. To read it in, we will use `read.csv`, which tells R to expect a comma-delimited file. Oh, and we should also get rid of those nasty factors for now.

```
aerial=read.csv(
  "https://raw.githubusercontent.com/cfrost3/MBM_R_Short_Course/master/BLSC_2018_RawObs_Fake.csv",
  header=TRUE,
  stringsAsFactors=FALSE)

str(aerial)
```

```
## 'data.frame':    624 obs. of  16 variables:
## $ Year      : int  2018 2018 2018 2018 2018 2018 2018 2018 2018 2018 ...
## $ Month     : chr   "6" "6" "6" "6" ...
## $ Day       : int   15 15 15 15 15 15 15 15 15 15 ...
## $ Seat      : chr   "LF" "LF" "LF" "LF" ...
## $ Observer  : chr   "FAKE" "FAKE" "FAKE" "FAKE" ...
## $ Strata    : chr   NA NA NA NA ...
## $ Transect  : chr   "19" "19" "19" "19" ...
## $ Segment   : logi  NA NA NA NA NA NA ...
## $ Flight_Dir: chr   "88" "88" "89" "89" ...
## $ Wind_Dir  : chr   "45" "45" "45" "45" ...
## $ Wind_Vel  : chr   "24" "24" "24" "24" ...
## $ Lat       : num   61.1 61.1 61.2 61.2 61.2 ...
## $ Long      : num  -166 -165 -165 -165 -165 ...
## $ Species   : chr   "START" "SCAU4" "LTDUF4" "BLSC4" ...
## $ Num       : chr   "19" "1" "1" "1" ...
## $ Obs_Type  : chr   "open" "single" "single" "pair" ...
```

Hey! We have real (fake) data! Anything stand out immediately in the structure of the file?

2.3.1 Data Types

R has 4 basic data types you will use (6 in total, actually, including raw and complex):

1. Integer (1, 2, 3, 4)
2. Numeric (2, 4, 2.4, 45.48)
3. Character ("cat", "DOG", "15d", "and so on")
4. Logical (TRUE, FALSE)

2.3.2 Data Structures

Depending on how you combine data types, you can end up with one of several data structures:

1. Vector - a collection of all one data type
2. Matrix - multidimensional collection of vectors
3. Factor - nominal set of unique values and a vector of integer indices
4. List - open format to contain data structure elements
5. Data frame - our most common, discussed above

There are many other data structures in R, including user-defined and package-defined structures. We will mostly stick to the common ones.

2.3.3 Useful Data Structure Functions

Base R comes with many quick ways to assess your data structures. We will go into more depth later, but for now, explore commands such as:

```
a = c(1,2,3,4,10.1)
```

```
length(a)  #how long is the vector?
```

```
## [1] 5
```

```
mean(a)    #mean value
```

```
## [1] 4.02
```

```
var(a)     #variance
```

```
## [1] 12.802
```

```
typeof(a)  #what data type? (double means floating decimal numeric)
```

```
## [1] "double"
```

```
is.na(a)   #any NA or missing values?
```

```
## [1] FALSE FALSE FALSE FALSE FALSE
```



```
summary(a) #some basic summarizing statistics
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00    2.00    3.00    4.02    4.00   10.10
```

```
max(a) #maximum value
```

```
## [1] 10.1
```

```
min(a) #minimum value
```

```
## [1] 1
```

```
dim(small.data) #what are the dimensions of a structure?
```

```
## [1] 4 5
```

```
class(small.data) #what class of structure is it?
```

```
## [1] "data.frame"
```

```
names(small.data) #what are the names of the objects that make up the structure?
```

```
## [1] "name"      "size"      "tenure"    "greatness" "interests"
```

```
unique(small.data$size) #what are all of the unique values in a range?
```

```
## [1] "small" "medium" "large"
```

3.0.0 Tidy Data

Tidy data refers to a data set that satisfies 3 conditions:

1. Every column is 1 variable
2. Every row is 1 observation
3. Every cell is 1 value

In R, tidy data lends itself seamlessly to the vectorized functions we discussed above (and thousands of others). In addition to summarizing functions, such as `length()`, `mean()`, `max()`, `min()`, etc., plot functions in R work sometimes exclusively and other times much cleaner and more efficiently with tidy data.

It is important to note here that quality controlled data (typos removed, missing data correctly recorded, etc.) and tidy data are 2 separate and important characteristics.

Messy data (data that is not tidy) is often the result of database design by a data collector, rather than an analyst or data manager. This is not meant to place blame. A data collector has one primary purpose; translate the raw observations in the field into some sort of tabular archive. That archive will generally, then, take the form and function that was easiest and fastest for the collector to enter the data (a very reasonable result since it is often preferable to get data collected and archived quickly). A data analyst similarly can have one primary purpose; to analyze data. While the analyst would also like to do things relatively as quickly as possible, messy data can make analysis lengthy, difficult, costly, and sometimes impossible. In MBM, for example, we have spent easily over 90% of our time cleaning data instead of on analysis (and the general consensus among similar analysts across all fields is 75-80% of time on cleaning).

3.1.0 Common Problems

There are 2 common problems with data that we run into frequently in MBM:

1. Using values as column names
2. Multiple variables in one column

3.1.1 Values as Column Names

Consider the following data frame:

```
messy1 = data.frame(Species = rep(c("COGO", "BOOW"), each = 3),
                    Box=c(1:6),
                    Visit1 = c(5,6,6,4,5,4),
                    Visit2 = c(5,6,5,4,5,7),
                    Visit3 = c(0,6,5,4,5,0),
                    Visit4 = c(NA,0,0,1,0,NA),
                    Visit5 = c(NA,NA,NA,0,NA,NA),
                    stringsAsFactors = FALSE
                    )
```

messy1

##	Species	Box	Visit1	Visit2	Visit3	Visit4	Visit5
## 1	COGO	1	5	5	0	NA	NA
## 2	COGO	2	6	6	6	0	NA
## 3	COGO	3	6	5	5	0	NA
## 4	BOOW	4	4	4	4	1	0
## 5	BOOW	5	5	5	5	0	NA
## 6	BOOW	6	4	7	0	NA	NA

Imagine they are visits to nests of 2 different species and the counts of the eggs that were found on each visit. Why is this messy? Looks clean enough. Easy enough for the data collector. Notice how columns 3-7 are labeled? This is a prime example of using a value for a variable (in this case, visit number) as the column header. It now becomes awkward for those nests that don't share the same number of visits. As an analyst, how can I script a quick summary of the number of times a nest was visited? Maybe by counting the number of cells that have a value greater than or equal to 0 in a row?

```
#give me the total of row 1 values greater than or equal to 0
```

```
v = sum(messy1[1,]>0)
```

v

```
## [1] NA
```

Ouch. Oh yeah, the NA values. A sum containing an NA is always NA. So let's remove those.

```
#give me the total of row 1 values greater than or equal to 0, and skip the NAs
```

```
v = sum(messy1[1,]>0, na.rm = TRUE)
```

v

```
## [1] 4
```

Ok, looks like it worked. But not when we spot check it. We forgot that Box will (probably) always be greater than or equal to 0. So let's tack on a modifier.

```
#give the total of row 1 values >= to 0, skip the NAs, subtract 1 for Box
v = sum(messy1[1,]>0, na.rm = TRUE)-1

v
```

```
## [1] 3
```

Well, we now correctly show 3 visits. But what could possibly go wrong with this in a larger script? Or imagine in an ACTUAL data set, where we have dozens of columns. This is the kind of thing that would drive an analyst bonkers AND would be a nightmare to fix after it has been done for years on a large data set. We won't even attempt to reshape this using base R commands. Lucky for us, we have the tidyverse package.

Tidyverse is actually a collection of packages developed by data scientists, for data scientists. When you install tidyverse, you get them all. It streamlines importing, cleaning, tidying, visualizing, analyzing, and reporting any size data set. Most (all) of the functions and functionality available in the tidyverse are also possible in base R, but often with significantly more (and generally more confusing) code. We can't possibly describe them all, but we will hit some of the main functions.

First, let's install the tidyverse package. This may take a while depending on how many of the packages you already have. We will also reload our fake aerial survey data to play with.

```
install.packages("tidyverse")
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 3.6.3
```

```
## Warning: package 'ggplot2' was built under R version 3.6.3
```

```
## Warning: package 'tibble' was built under R version 3.6.3
```

```
## Warning: package 'tidyr' was built under R version 3.6.3
```

```
## Warning: package 'readr' was built under R version 3.6.1
```

```
## Warning: package 'purrr' was built under R version 3.6.3
```

```
## Warning: package 'dplyr' was built under R version 3.6.3
```

```
## Warning: package 'forcats' was built under R version 3.6.3
```

```
aerial=read.csv(
  "https://raw.githubusercontent.com/cfrost3/MBM_R_Short_Course/master/BLSC_2018_RawObs_Fake.csv",
  header=TRUE,
  stringsAsFactors=FALSE)
```

If you recall, we used the \$ operator to reference a column in a data frame. Tidyverse provides a cleaner syntax with the select function.

```
#grab Transect, Species, and Num only from aerial
aerial.subset = select(aerial, Transect, Species, Num)

head(aerial.subset) #only show me the first few rows
```

```
##   Transect Species Num
## 1      19   START  19
## 2      19   SCAU4   1
## 3      19  LTDUF4   1
## 4      19   BLSC4   1
## 5      19   SCAU2   1
## 6      19   BLSC4   1
```

Quick, easy, and intuitive. Now let's only keep observations of 3 or more.

```
aerial.subset2 = filter(aerial.subset, Num > 2) #give me observations greater than 2

head(aerial.subset2)
```

```
##   Transect Species Num
## 1      15   SCAU1   5
## 2      16   BLSC3   3
## 3      27   START  27
## 4      27  ENDPT  27
## 5      26   START  26
## 6      26  BLSCF2   4
```

These are useful, but notice how we keep having to save out our results as new names with an increasing number on the end? This will get messy and out of control fast as we add on more filters and subsets and summaries. To combat this, the Tidyverse (specifically the magrittr package) has the pipe operator (%>%). The pipe operator chains together a series of functions and uses the results of the previous pipes in the subsequent pipes. It makes for a cleaner sequence that you can track later. For example, the sequence below will subset our columns and filter to bigger observations as we did above, but also filter out the START and ENDPT observations, then count the number of observations by species that met the conditions leading up to the aggregate.

```
obs.by.species = aerial %>%

  #grab Transect, Species, and Num only from aerial
  select(Transect, Species, Num) %>%

  #Num greater than 2, but remove START and ENDPT observations
  filter(Num > 2 & !(Species %in% c("START", "ENDPT"))) %>%

  aggregate(Num~Species, ., FUN=length)

  names(obs.by.species)[2]="N.obs"

obs.by.species
```

```
##   Species N.obs
```

```
## 1 BLSC3 2
## 2 BLSC4 1
## 3 BLSCF2 1
## 4 BLSCF3 2
## 5 LTDU3 1
## 6 SCAU1 6
## 7 SCAU2 2
## 8 SCAU3 7
```

We have a lot going on there. Note the command `%in%` that will only keep string values that it finds within a set of other string values. But in this case, I wrap it in `!`, which can be read as “not.” So that line becomes “filter down to observations greater than 2 and where Species is not START or ENDPT.”

Now that we have the pipe in our toolbox, let’s revisit the actual messy data problem.

```
messy1
```

```
## Species Box Visit1 Visit2 Visit3 Visit4 Visit5
## 1 COGO 1 5 5 0 NA NA
## 2 COGO 2 6 6 6 0 NA
## 3 COGO 3 6 5 5 0 NA
## 4 BOOW 4 4 4 4 1 0
## 5 BOOW 5 5 5 5 0 NA
## 6 BOOW 6 4 7 0 NA NA
```

We have 6 nest boxes that were visited up to 5 times and we just want a summary of the number of visits and the average number of eggs in the nest each visit. We have incorrectly recorded values of our variable Visit as column headers. Let’s fix it.

```
tidy1 = messy1 %>%
  pivot_longer(c("Visit1", "Visit2", "Visit3", "Visit4", "Visit5",),
    names_to = "Visit",
    values_to = "eggs")
tidy1
```

```
## # A tibble: 30 x 4
## Species Box Visit eggs
## <chr> <int> <chr> <dbl>
## 1 COGO 1 Visit1 5
## 2 COGO 1 Visit2 5
## 3 COGO 1 Visit3 0
## 4 COGO 1 Visit4 NA
## 5 COGO 1 Visit5 NA
## 6 COGO 2 Visit1 6
## 7 COGO 2 Visit2 6
## 8 COGO 2 Visit3 6
## 9 COGO 2 Visit4 0
## 10 COGO 2 Visit5 NA
## # ... with 20 more rows
```

This result is...ehhhh. It looks ok, but we can do better.

```
tidy1 = messy1 %>%
  pivot_longer(
    cols = starts_with("Visit"),
    names_to = "Visit",
    names_prefix = "Visit",
    values_to = "Eggs",
    values_drop_na = TRUE
  )

head(tidy1,10)
```

```
## # A tibble: 10 x 4
##   Species Box Visit Eggs
##   <chr>   <int> <chr> <dbl>
## 1 COGO    1 1      5
## 2 COGO    1 2      5
## 3 COGO    1 3      0
## 4 COGO    2 1      6
## 5 COGO    2 2      6
## 6 COGO    2 3      6
## 7 COGO    2 4      0
## 8 COGO    3 1      6
## 9 COGO    3 2      5
## 10 COGO   3 3      5
```

We took messy1 and applied pivot_longer, which will translate our column values into expanded row values. We further parameterized it by asking for all columns that started with Visit (cols = starts_with("Visit")), and removed the prefix Visit (names_prefix = "Visit") and sent the remainder to a new Visit column (names_to = "Visit"). We then took the row values (counts of eggs, though you can't tell from the data) and sent them to the appropriate new Species-Box-Visit row combination and labelled the new column Eggs (values_to = "Eggs"). Finally, we removed the NAs. These were causing problems above and since they were really just empty cells and not visits, they are better off gone. Now we can cleanly summarize our data!

```
vis = tidy1 %>%
  aggregate(Visit~Species + Box, ., FUN=max)
names(vis)[3] = "Num.Visits"

vis
```

```
##   Species Box Num.Visits
## 1   COGO   1         3
## 2   COGO   2         4
## 3   COGO   3         4
## 4   BOOW   4         5
## 5   BOOW   5         4
## 6   BOOW   6         3
```

```
avg.egg = tidy1 %>%
  filter(Visit == 1) %>%
  aggregate(Eggs~Species, ., FUN=mean)
names(avg.egg)[2] = "Avg.Eggs"

avg.egg
```

```
## Species Avg.Eggs
## 1 BOOW 4.333333
## 2 COGO 5.666667
```

#Or using just tidyverse

```
tidy1 %>%
  filter(Visit == 1) %>%
  group_by(Species) %>%
  summarize(Avg.Eggs=mean(Eggs))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 2 x 2
## Species Avg.Eggs
## <chr>      <dbl>
## 1 BOOW      4.33
## 2 COGO      5.67
```

3.1.2 Multiple Variables in One Column

Another common problem in MBM data sets is combining the values of multiple variables into a single column. Sometimes this is what we really thought was the best way to store values, and other times we make excuses about the limitations of our data collection programs or devices to store things certain ways. The results are messy data either way, and this serves to increase the rate of decay of the actual information we can extract from our data. Consider the following (fake) data:

```
messy2 = data.frame(monthday=paste(rep(6,10), c(1:10), sep="/"), Grade=rep(c("A", "F"), 5))
```

```
messy2
```

```
## monthday Grade
## 1 6/1 A
## 2 6/2 F
## 3 6/3 A
## 4 6/4 F
## 5 6/5 A
## 6 6/6 F
## 7 6/7 A
## 8 6/8 F
## 9 6/9 A
## 10 6/10 F
```

In this simple example, we have monthday, which is the month and day separated by a period, and a fictitious grade received on that day. We should really have month and day as their own columns. Let's separate.

```
messy2 %>%
  separate(monthday, into=c("Month", "Day"), sep="/")
```

```
## Month Day Grade
## 1 6 1 A
```

```
## 2      6      2      F
## 3      6      3      A
## 4      6      4      F
## 5      6      5      A
## 6      6      6      F
## 7      6      7      A
## 8      6      8      F
## 9      6      9      A
## 10     6     10     F
```

That was suspiciously easy...so let's try it with our aerial data. The column Species sounds like it should be the species. But it isn't. It is actually 3 values: species, behavior, and distance. But that isn't the end of it. The behavior was only recorded if the bird was flying. Pretty messy. So let's tidy it.

```
unique(aerial$Species)
```

```
## [1] "START" "SCAU4" "LTUDF4" "BLSC4" "SCAU2" "BLSCF3" "BLSC3"
## [8] "BLSCF4" "SCAU3" "BLSC2" "LTDU4" "SCAU1" "LTDU3" "BLSC1"
## [15] "ENDPT" "LTDU2" "LTDU1" "BLSCF1" "BLSCF2" "WWSC2"
```

```
aerial.tidy = aerial %>%
  filter(!(Species %in% c("START", "ENDPT"))) %>% #remove start and end points
  select(Lat, Long, Species, Num) %>% #filter to just 4 columns for illustration
  separate(Species, into=c("Species", "Distance"), sep = -1) %>% #take away the rightmost string value
  separate(Species, into=c("Species", "Behavior"), sep = 4) #take away the first 4 string values

head(aerial.tidy)
```

```
##      Lat      Long Species Behavior Distance Num
## 1 61.1457 -165.4586   SCAU          4      1
## 2 61.1564 -164.8116   LTDU          F      4      1
## 3 61.1587 -164.6340   BLSC          4      1
## 4 61.1595 -164.5453   SCAU          2      1
## 5 61.1605 -164.4803   BLSC          4      1
## 6 61.1605 -164.4803   BLSC          4      1
```

```
unique(aerial.tidy$Species)
```

```
## [1] "SCAU" "LTDU" "BLSC" "WWSC"
```

An initial look at the Species column showed that the Distance was always the last value in the string. So we were able to tell separate that we need to pull that one out.

```
separate(Species, into=c("Species", "Distance"), sep = -1)
```

Then after pulling that one out, we saw that the species code was always the first 4 characters in the string, so we were able to tell separate to always give us those.


```
separate(Species, into=c("Species", "Behavior"), sep = 4)
```

With the pipe, we did it all in one command, and we can move on to bigger and better things, such as visualization!

4.0.0 Tables and Figures

The first step in any data manipulation, analysis, QA/QC process, or almost any other data-related activity should be to visualize your data. This will be the best way to catch any potential inconsistencies or errors before you embed them deep in an analysis or report. Base R provides many ways to visualize data quickly and easily, but contributed packages such as DT, kableExtra, ggplot2, and leaflet expand on base R to provide publication-quality tables and figures. We will start with base R tables.

4.1.1 Base R Tables

Let's make sure we have our aerial data loaded and tidy.

```
library(tidyverse)

aerial=read.csv(
  "https://raw.githubusercontent.com/cfrost3/MBM_R_Short_Course/master/BLSC_2018_RawObs_Fake.csv",
  header=TRUE,
  stringsAsFactors=FALSE)

aerial.tidy = aerial %>%
  filter(!(Species %in% c("START", "ENDPT"))) %>% #remove start and end points
  select(Lat, Long, Species, Num) %>% #filter to just 4 columns for illustration
  separate(Species, into=c("Species", "Distance"), sep = -1) %>% #take away the rightmost string value
  separate(Species, into=c("Species", "Behavior"), sep = 4) #take away the first 4 string values

head(aerial.tidy)
```

```
##           Lat      Long Species Behavior Distance Num
## 1 61.1457 -165.4586   SCAU              4      1
## 2 61.1564 -164.8116   LTDU              F      4      1
## 3 61.1587 -164.6340   BLSC              4      1
## 4 61.1595 -164.5453   SCAU              2      1
## 5 61.1605 -164.4803   BLSC              4      1
## 6 61.1605 -164.4803   BLSC              4      1
```

```
table(aerial.tidy$Species,aerial.tidy$Distance)
```

```
##
##           1      2      3      4
## BLSC    23    43   146    42
## LTDU      5      7    11      2
## SCAU    54    53   120    12
## WWSC      0      1      0      0
```

The format is `table(rows, columns)`, and we can even add multidimensionality:

```
table(aerial.tidy$Species, aerial.tidy$Distance, aerial.tidy$Behavior)
```

```
## , , =
##
##
##      1    2    3    4
## BLSC 21  42 136  39
## LTDU  5   7  11   1
## SCAU 54  53 120  12
## WWSC  0   1   0   0
##
## , , = F
##
##
##      1    2    3    4
## BLSC  2   1  10   3
## LTDU  0   0   0   1
## SCAU  0   0   0   0
## WWSC  0   0   0   0
```

One important thing to remember with `table` in base R is that the default summary will simply give you a count of the number of times the 2 observations occurred together (NOT the total number of things we counted by species in this example). Always keep track of what you are visualizing. You can also easily calculate the proportion of observations in a table.

```
prop.table(table(aerial.tidy$Species, aerial.tidy$Distance))
```

```
##
##      1          2          3          4
## BLSC 0.044315992 0.082851638 0.281310212 0.080924855
## LTDU 0.009633911 0.013487476 0.021194605 0.003853565
## SCAU 0.104046243 0.102119461 0.231213873 0.023121387
## WWSC 0.000000000 0.001926782 0.000000000 0.000000000
```

This proportion table tells us that our most frequent observation (roughly 28% of the time) was BLSC in distance bin 3. Notice how we had to wrap the `table()` function with `prop.table()`? That is because while `table()` expects (rows, columns) format, `prop.table` expects a `table()` as input. We can even wrap the whole thing in a `max()` to verify our 28%.

```
max(prop.table(table(aerial.tidy$Species, aerial.tidy$Distance)))
```

```
## [1] 0.2813102
```

One thing we would certainly want to know from our aerial data collection is the number (Num) of each species we counted. We can get that number (albeit inefficiently) using tables.

```
table(aerial.tidy$Species, aerial.tidy$Num)
```

```
##
##           1  10   2  25   3  30   4   5   6   7   d   t
##   BLSC 245   0   3   0   4   0   1   0   1   0   0   0
##   LTDU  24   0   0   0   0   0   0   0   0   0   0   1
##   SCAU 215   1   8   2   7   1   1   1   1   1   1   0
##   WWSC   1   0   0   0   0   0   0   0   0   0   0   0
```

```
str(aerial.tidy)
```

```
## 'data.frame':   519 obs. of  6 variables:
##  $ Lat      : num  61.1 61.2 61.2 61.2 61.2 ...
##  $ Long     : num -165 -165 -165 -165 -164 ...
##  $ Species  : chr  "SCAU" "LTDU" "BLSC" "SCAU" ...
##  $ Behavior : chr  "" "F" "" "" ...
##  $ Distance : chr  "4" "4" "4" "2" ...
##  $ Num      : chr  "1" "1" "1" "1" ...
```

Well, that won't work. It seems we've identified a QA/QC issue. We could've caught this by checking the structure of our data frame as well (see how Num shows as chr?). Let's pluck those 2 offenders out and fix them. We will find them and fix them in 2 different ways for illustration only. One useful function for this is which(). It will return the index or indices (similar to the row number in Excel) of the row(s) of the observation(s) that fit a condition. In our case, we want to know where the d and t are hanging out.

```
which(aerial.tidy$Num %in% c("d", "t")) #what positions are the d and t in?
```

```
## [1] 427 438
```

```
aerial.tidy$Num[427] #verify
```

```
## [1] "d"
```

```
aerial.tidy[427,] #show me that whole row 427 and all of its columns
```

```
##           Lat      Long Species Behavior Distance Num
## 427 66.2326 -165.6542   SCAU              1      d
```

```
aerial.tidy[c(427,438),] #show me both row 427 and 438 and all of their columns
```

```
##           Lat      Long Species Behavior Distance Num
## 427 66.2326 -165.6542   SCAU              1      d
## 438 66.2322 -164.0431   LTDU              3      t
```

Let's imagine that we had our observer go back through the recordings (because both the observer and recordings still exist), and they found that d should be 2 and t should be 1. Since we know the indices of the offenders, we can just overwrite them.

```
aerial.tidy$Num[427] = 2
```

Alternatively, we can grab the index and do something to it all at once:

```
aerial.tidy$Num[aerial.tidy$Num == "t"] = 1
```

#find any values of Num == "t" and replace with 1

If we table it again, we can see that we have at least removed the d and t observations. But the structure shows Num still contained as character instead of numeric.

```
table(aerial.tidy$Species,aerial.tidy$Num)
```

```
##
##      1  10   2  25   3  30   4   5   6   7
## BLSC 245   0   3   0   4   0   1   0   1   0
## LTDU  25   0   0   0   0   0   0   0   0   0
## SCAU 215   1   9   2   7   1   1   1   1   1
## WWSC   1   0   0   0   0   0   0   0   0   0
```

```
str(aerial.tidy)
```

```
## 'data.frame':   519 obs. of  6 variables:
## $ Lat      : num  61.1 61.2 61.2 61.2 61.2 ...
## $ Long     : num -165 -165 -165 -165 -164 ...
## $ Species  : chr  "SCAU" "LTDU" "BLSC" "SCAU" ...
## $ Behavior: chr   "" "F" "" "" ...
## $ Distance: chr  "4" "4" "4" "2" ...
## $ Num      : chr  "1" "1" "1" "1" ...
```

To fix that, we can just tell aerial.tidy to reconsider that column as numeric. If we get any warnings or errors, we can assume there were other offenders in there. In our case, there weren't. Now when we table it, the columns even reorder correctly.

```
aerial.tidy$Num = as.numeric(aerial.tidy$Num)
```

```
str(aerial.tidy)
```

```
## 'data.frame':   519 obs. of  6 variables:
## $ Lat      : num  61.1 61.2 61.2 61.2 61.2 ...
## $ Long     : num -165 -165 -165 -165 -164 ...
## $ Species  : chr  "SCAU" "LTDU" "BLSC" "SCAU" ...
## $ Behavior: chr   "" "F" "" "" ...
## $ Distance: chr  "4" "4" "4" "2" ...
## $ Num      : num   1 1 1 1 1 1 1 1 1 1 ...
```

```
table(aerial.tidy$Species,aerial.tidy$Num)
```

```
##
##      1   2   3   4   5   6   7  10  25  30
## BLSC 245   3   4   1   0   1   0   0   0   0
## LTDU  25   0   0   0   0   0   0   0   0   0
## SCAU 215   9   7   1   1   1   1   1   2   1
## WWSC   1   0   0   0   0   0   0   0   0   0
```

Now if we want a basic summary of our observations, we can use `margin.table` to sum the table across either or both dimensions. In other words, a sum by rows, a sum by columns, or a total sum (keeping in mind our unit is observation and NOT number of things seen). Note that `margin.table` is like `prop.table` and the argument passed must be a table object.

```
#The 1 specifies the first (row) dimension (how many times did we observe BLSC?)
margin.table(table(aerial.tidy$Species,aerial.tidy$Num), 1)
```

```
##
## BLSC LTDU SCAU WWSC
## 254 25 239 1
```

```
#The 2 specifies the second (column) dimension (how many times did we see only 1 of something?)
margin.table(table(aerial.tidy$Species,aerial.tidy$Num), 2)
```

```
##
## 1 2 3 4 5 6 7 10 25 30
## 486 12 11 2 1 2 1 1 2 1
```

```
#No dimension argument gives us the sum over both rows and columns (how many total observations were there?)
margin.table(table(aerial.tidy$Species,aerial.tidy$Num))
```

```
## [1] 519
```

4.1.2 Tables in Contributed Packages

We will demonstrate table functions from 2 contributed (user-developed) packages, `DT` and `kableExtra`. `DT` stands for data table and works best to visualize large data sets in a clean, searchable, sortable, filterable way. `kableExtra` is a handy package for creating tables that use the functionality of html coding to render in either html or pdf reports. We need to make sure we have both installed and loaded.

```
install.packages(c("DT", "kableExtra", "htmltools", "zoo"))
```

```
library(DT)
library(kableExtra)
library(AKAerial)
library(tidyverse)
library(htmltools)
library(zoo)
```

For the first example, we will load a table of real aerial index estimates from the `AKAerial` package and use some of our tidyverse functions to trim it down for a specific table using `datatable` from the package `DT`.

```
ACPHistoric$combined %>%
  filter(Species=="SPEI") %>%
  select(Year, itotal, itotal.se) %>%
  datatable()
```

Show 10 entries

Search:

	Year	Itotal	Itotal.se
1	2007	4816.20682397878	697.895902107711
2	2008	5797.65876644687	676.275377587682
3	2009	5224.29377755009	1041.10982461076
4	2010	5663.48541997397	820.266858935159
5	2011	7867.67838345973	956.368696239887
6	2012	4631.63471765462	460.50390706327
7	2013	7433.2610355183	1106.55275752308
8	2014	7103.33028723992	1251.23339858375
9	2015	5735.91785052616	796.900972060308
10	2016	4431.0195685689	881.213183006097

Showing 1 to 10 of 13 entries

Previous 1 2 Next

That's a solid first effort, but we can do better. Our itotal column really shouldn't allow decimal amounts of birds, and the standard error could use some trimming. I also don't particularly like seeing the row numbers in there. And the column names could be better. And we need a title, maybe shrink the width, center that middle column, and show me all of the entries at once since there are only 13.

```
ACPHistoric$combined %>%
  filter(Species=="SPEI") %>%
  select(Year, itotal, itotal.se) %>%
  mutate_at("itotal", ~round(., 0)) %>%
  mutate_at("itotal.se", ~round(., 2)) %>%
  datatable(rownames=FALSE, #cut out those row numbers showing
            fillContainer = TRUE, #auto-size the column widths
            colnames=c("Year","Indicated Total","SE"), #change column names
            #set the caption obtusely using html styling
            caption = htmltools::tags$caption( style = 'caption-side: top;
                                                text-align: center;
                                                color:black;
                                                font-size:100% ;',
                                                'Table 1: Indicated Total SPEI on the ACP, 2007-2019.')
```

```
options=list(
  autoWidth=TRUE,
  #center the column indexed 1 (starts with 0)
  columnDefs = list(list(className = 'dt-center', targets = 1)),
  pageLength=length(.[,1]), #display all of the data points
  dom="") #take away the search and page functionality
```

Table 1: Indicated Total SPEI on the ACP, 2007-2019.

Year	Indicated Total	SE
2007	4816	697.9
2008	5798	676.28
2009	5224	1041.11
2010	5663	820.27
2011	7868	956.37
2012	4632	460.5
2013	7433	1106.55
2014	7103	1251.23
2015	5736	796.9
2016	4431	881.21
2017	4465	751.65
2018	4761	1052.87
2019	3754	567.06

Now that's a table! It renders a little small though in this pdf since html styling doesn't jive well with pdf output. It should look great in your console or in an html report. We used a new tidyverse function `mutate_at`. That changes the styling or content of certain columns only, referenced by name or number. In this case, we applied the function `round()` to the piped data frame and told it to use 0 and 2 decimal places. The rest of the changes were implemented as arguments to `datatable()`.

Now we will implement a similar table using `kableExtra`.

```
YKGHistoric$combined %>%
  filter(Species=="BRAN") %>%
  select(Year, itotal, itotal.se) %>%
  kable()
```

Year	itotal	itotal.se
1985	4587.359	1377.462
1986	13065.045	2009.412
1987	12316.825	1228.657
1988	19774.425	2354.796
1989	26414.090	3239.270
1990	28371.901	3450.564
1991	21219.565	2389.966
1992	19531.009	1910.923
1993	31543.300	3336.424
1994	30487.396	2591.968
1995	34041.643	3187.085
1996	29077.673	2365.832
1997	30450.473	2757.902
1998	20518.538	1500.905
1999	21101.275	1846.622
2000	24619.576	2270.803
2001	30592.181	2658.417
2002	19623.449	1361.620
2003	20019.753	1555.904
2004	19172.114	1466.754
2005	20666.951	1734.801
2006	18835.411	1616.762
2007	19891.241	1858.726
2008	28361.608	1930.942
2009	23000.566	1516.937
2010	23107.933	1811.589
2011	16261.917	1243.417
2012	21650.089	2671.557
2013	23867.171	2953.542
2014	23160.193	2394.913
2015	20856.023	1532.706
2016	29738.491	2331.881
2017	21038.565	1796.098
2018	26132.308	2317.572
2019	22099.878	1869.757

That might be good enough for some journals...but not for me. Let's get our rounding going, change those column names, and even add a 3-year and 10-year rolling average (from the zoo package). Eh...and a caption, a footnote, and a color scheme, just to give it that "next-level" feel.

```
YKGHistoric$combined %>%
  filter(Species=="EMGO") %>%

  select(Year, itotal, itotal.se) %>%

  #rollapply in zoo rolls through a vector and applies a function to the segments

  mutate(avg3=rollapply(itotal,3,mean,align='right',fill=NA)) %>%

  mutate(avg10=rollapply(itotal,10,mean,align='right',fill=NA)) %>%

  mutate_at(c("itotal", "avg3", "avg10"), ~round(., 0)) %>%
```



```
mutate_at("itotal.se", ~round(., 2)) %>%

#apply our conditional population objective coloring
mutate_at("itotal", function(x) {
  cell_spec(x, "html", color = ifelse(x > 26000, "green", "red"))
}) %>%

kable(format="html",
  escape = F, #html scheme requirement to make the color statement work
  #now we paste in the footnote denotation on the indicated total column
  col.names = c("Year", "Indicated Total*",
    "SE", "3-year Avg", "10-year Avg"),
  #and add the top caption
  caption = "Table 2: Indicated Total EMGO on the YK-Delta, 1985-2019, including 3- and 10-year A

#adding the footnote is done in its own function so we pipe the kable to it

footnote(symbol= "Indicated Total is used by the AMBCC to regulate harvest. Values of 26000 and above

#this tells us to use the default bordered style
kable_styling("bordered",
  full_width=FALSE,
  font_size = 14)
```

Table 2: Indicated Total EMGO on the YK-Delta, 1985-2019, including 3- and 10-year Averages.

Year	Indicated Total*	SE	3-year Avg	10-year Avg
1985	18687	1572.22	NA	NA
1986	11355	707.81	NA	NA
1987	10612	886.73	13551	NA
1988	13175	813.72	11714	NA
1989	14340	829.85	12709	NA
1990	14609	859.16	14041	NA
1991	12432	969.35	13793	NA
1992	13251	700.33	13430	NA
1993	15524	993.62	13736	NA
1994	17101	834.46	15292	14108
1995	17463	863.03	16696	13986
1996	23578	2283.20	19381	15208
1997	22525	1294.91	21189	16400
1998	19714	1031.07	21939	17054
1999	20274	1180.34	20838	17647
2000	17260	698.78	19083	17912
2001	27674	1168.25	21736	19436
2002	19193	995.96	21376	20031
2003	20899	1313.53	22589	20568
2004	21514	831.30	20535	21009
2005	20739	1079.30	21051	21337
2006	26325	1346.79	22859	21612
2007	26281	1592.80	24448	21987
2008	22543	922.22	25050	22270
2009	20498	836.23	23107	22293
2010	19903	904.22	20981	22557
2011	21319	999.51	20573	21921
2012	20628	1324.44	20617	22065
2013	29876	1829.83	23941	22963
2014	31669	2662.58	27391	23978
2015	28634	1401.63	30060	24768
2016	34242	2004.61	31515	25559
2017	30090	1448.18	30988	25940
2018	30218	1527.67	31516	26708
2019	26583	1230.26	28963	27316

* Indicated Total is used by the AMBCC to regulate harvest. Values of 26000 and above result in an open harvest.

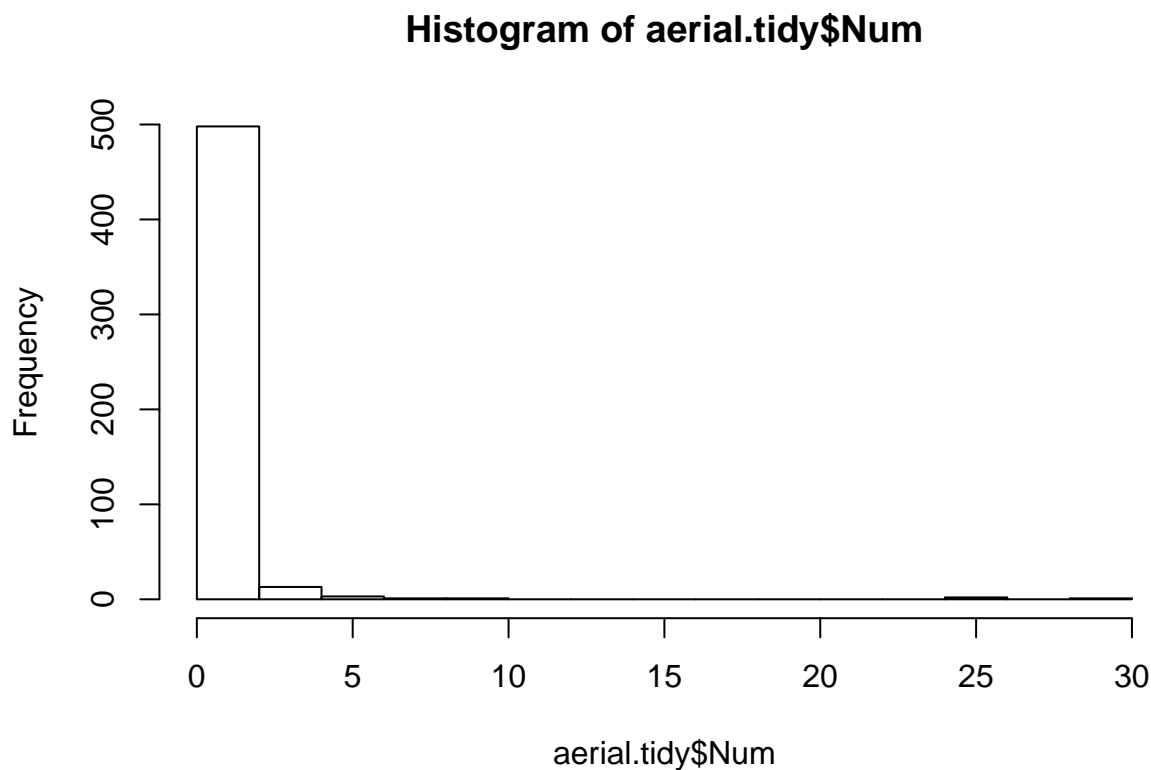
The result could easily be hung on the nearest refrigerator (sorry, the styling will look funny as a pdf). The majority of the work in creating a nice table is getting the formatting just right for your target. It will take some practice to figure out where the option is that changes each of your style elements. Don't forget that a package probably already exists to do exactly whatever manipulation you are dreaming of doing to your data. Search engines are powerful tools.

4.2.1 Plotting in Base R

Base R comes loaded with relatively robust plotting capability. Just like with tables, though, the R community identified many features that were either missing, or difficult to implement, and expanded the R universe with contributed plotting content. We will start in base R with the basics.

We will revisit our `aerial.tidy` that we created earlier. Note that we aren't applying any pipes just yet since these visualizations are usually 1 or 2 contrasts and nothing else gets displayed, unlike a table where we want to choose exactly what data doesn't get displayed.

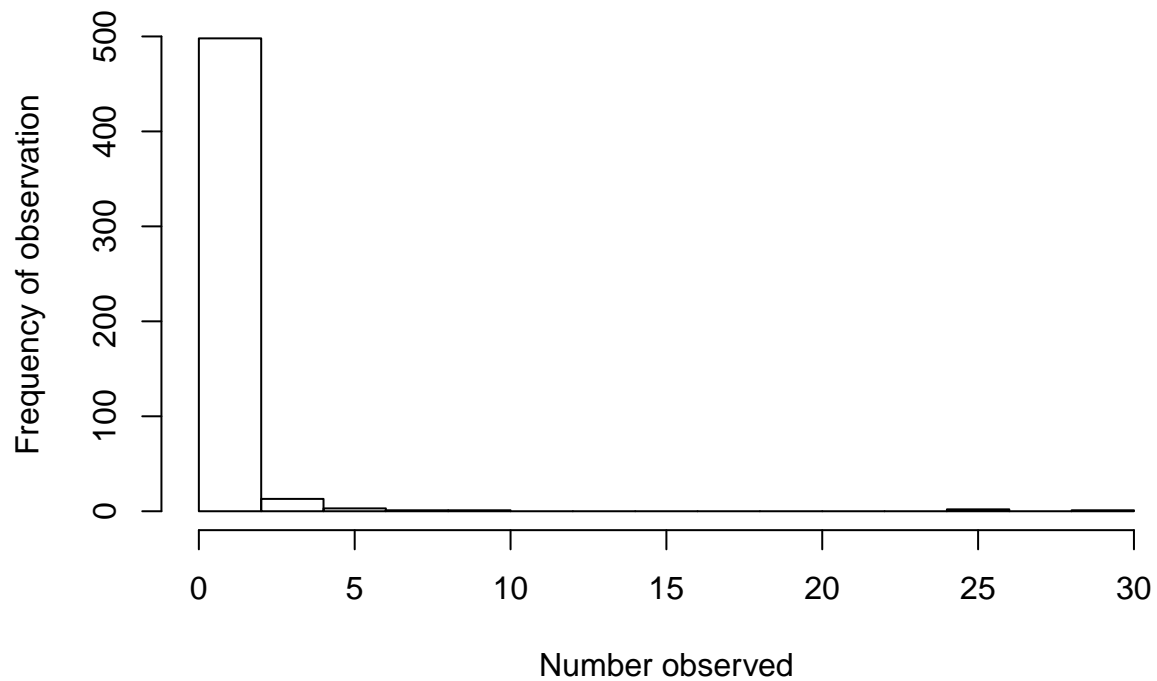
```
hist(aerial.tidy$Num)
```



The default x and y axis labels and main title leave us feeling a little empty. Let's fix.

```
hist(aerial.tidy$Num,  
     xlab = "Number observed",  
     ylab = "Frequency of observation",  
     main = "Observations by size and frequency")
```

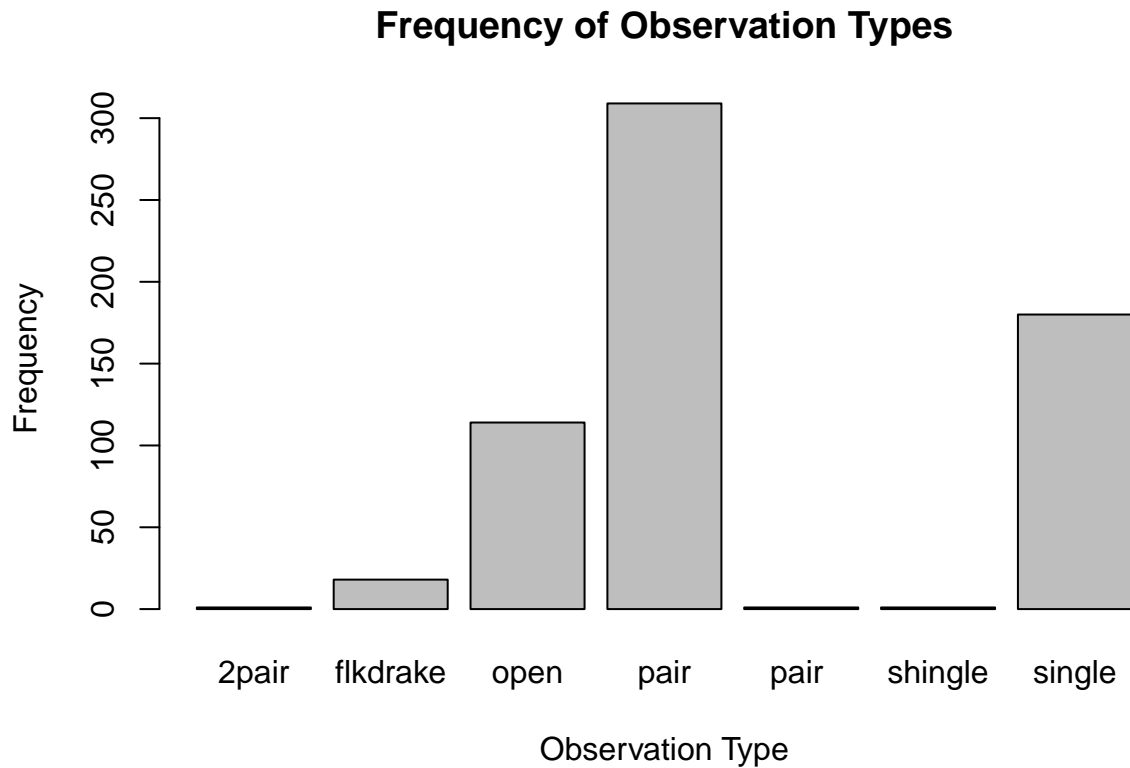
Observations by size and frequency



Now let's make a barplot. It looks similar to the histogram, but the histogram only works with numeric vectors. A barplot can work with strings and factors converted to tables.

```
counts=table(aerial$Obs_Type)

barplot(counts,
        xlab="Observation Type",
        ylab="Frequency",
        main="Frequency of Observation Types")
```



Hmmm... looks like more QA/QC problems. They seem pretty intuitive to fix except for one. The second case of “pair” on the x axis is actually “pair ” (with a space on the end). This is a common problem with character data and one that R doesn’t automatically know how to treat. We would have identified it using a `unique(aerial$Obs_Type)` call. So let’s fix these.

```
aerial$Obs_Type[aerial$Obs_Type %in% c("2pair", "pair ")] = "pair"

aerial$Obs_Type[aerial$Obs_Type == "shingle"] = "single"

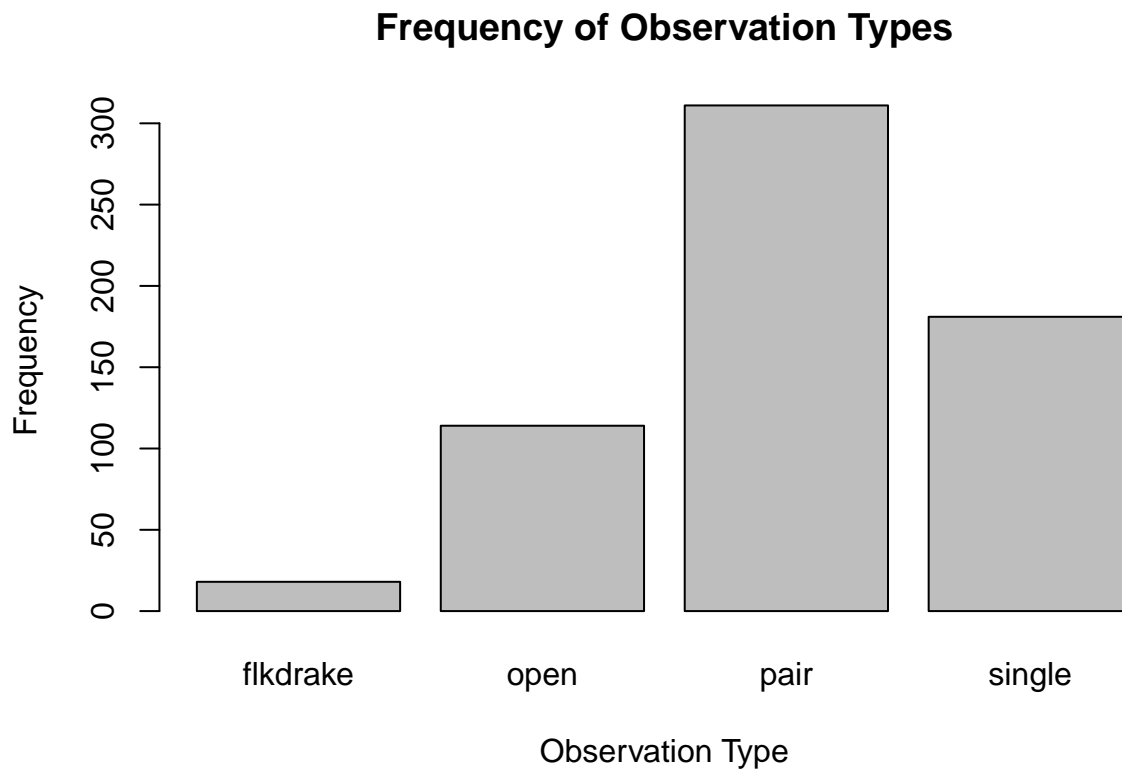
unique(aerial$Obs_Type)
```

```
## [1] "open"      "single"    "pair"      "flkdrake"
```

Ok, those look reasonable now. Again with the barplot.

```
counts=table(aerial$Obs_Type)


barplot(counts,
        xlab="Observation Type",
        ylab="Frequency",
        main="Frequency of Observation Types")
```



Looks better. We can even do 2-dimensional tables.

```
counts=table(aerial.tidy$Distance, aerial.tidy$Species)

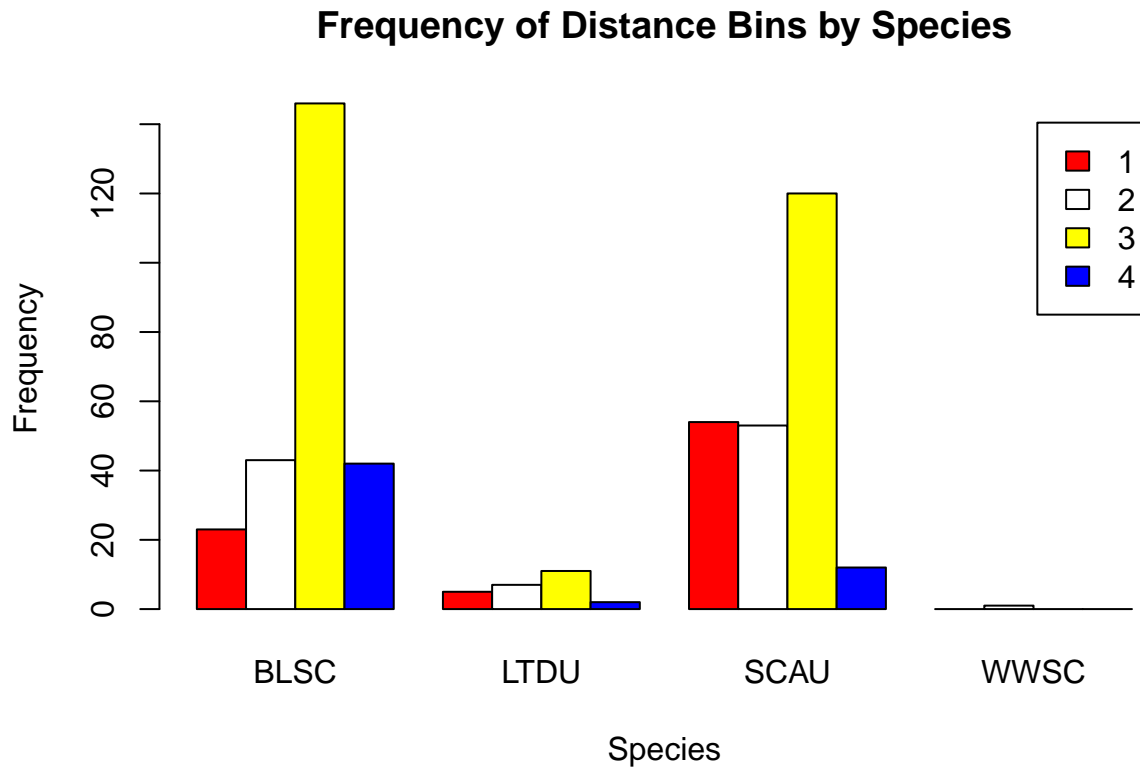
barplot(counts,
        xlab="Species",
        ylab="Frequency",
        main="Frequency of Distance Bins by Species",
        legend=rownames(counts),
        col=c("red","white","yellow","blue"))
```



You can place them side by side with the argument `beside = TRUE`.

```
counts=table(aerial.tidy$Distance, aerial.tidy$Species)

barplot(counts,
  xlab="Species",
  ylab="Frequency",
  main="Frequency of Distance Bins by Species",
  legend=rownames(counts),
  col=c("red","white","yellow","blue"),
  beside=TRUE)
```



You can even configure the R plot window to show you multiple figures at once by changing up the graphical parameters.

```
par(mfrow = c(2,2)) #give me a 2x2 matrix graphical display

#R will fill the 4 spaces with my next 4 plot calls

hist(aerial.tidy$Num,
     xlab = "Number observed",
     ylab = "Frequency of observation",
     main = "Observations by size and frequency")

counts=table(aerial$Obs_Type)

barplot(counts,
        xlab="Observation Type",
        ylab="Frequency",
        main="Frequency of Observation Types")

counts=table(aerial.tidy$Distance, aerial.tidy$Species)

barplot(counts,
        xlab="Species",
        ylab="Frequency",
```

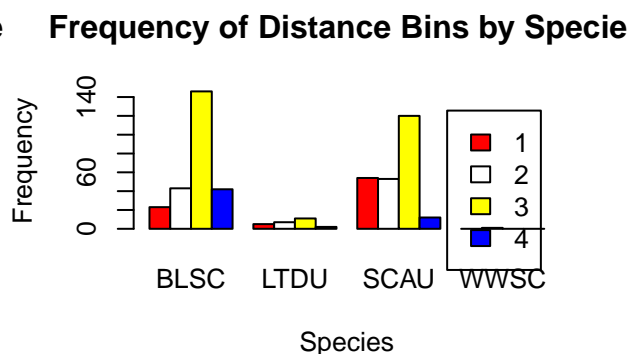
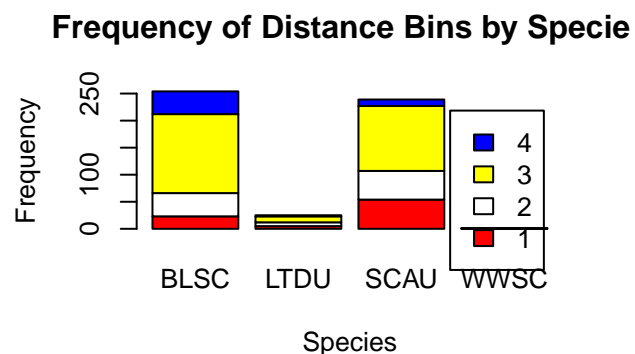
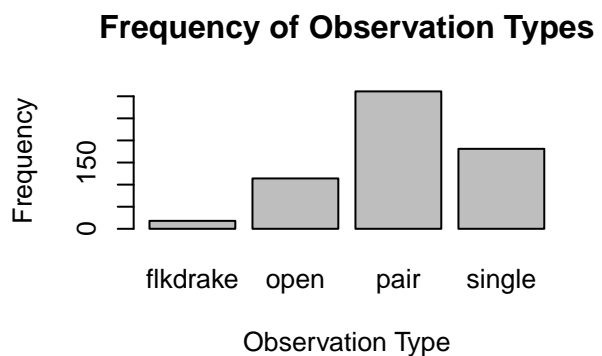
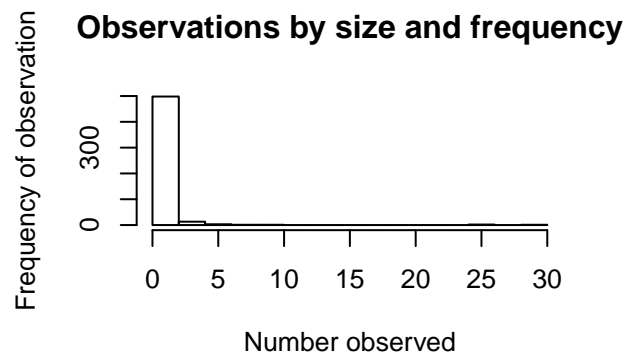
```

main="Frequency of Distance Bins by Species",
legend=rownames(counts),
col=c("red", "white", "yellow", "blue"))

counts=table(aerial.tidy$Distance, aerial.tidy$Species)

barplot(counts,
        xlab="Species",
        ylab="Frequency",
        main="Frequency of Distance Bins by Species",
        legend=rownames(counts),
        col=c("red", "white", "yellow", "blue"),
        beside=TRUE)

```



```

#don't forget to change it back!
par(mfrow=c(1,1))

```

The base functions `plot()`, `points()`, and `lines()` work splendidly for time series type data.

```

ltdu = YKDHistoric$combined %>%
  filter(Species == "LTDU")

plot(total~Year, data = ltdu,

```



```

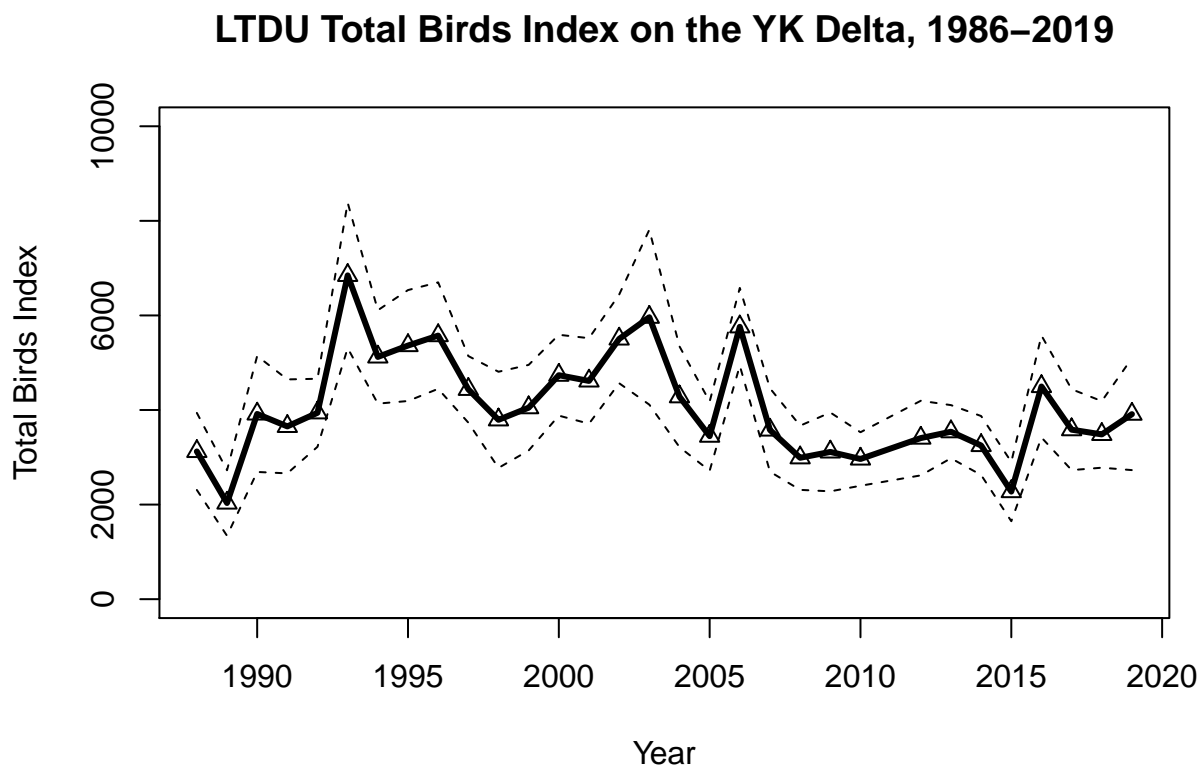
type = 'l', #connect the dots
lwd = 3, #thicken the line
ylim = c(0,10000), #define the y range instead of letting R calculate it
ylab = "Total Birds Index",
main = "LTDU Total Birds Index on the YK Delta, 1986-2019")

lines((total-1.96*total.se)~Year, data = ltdu, lty = 2)

lines((total+1.96*total.se)~Year, data = ltdu, lty = 2)

points(total~Year, data = ltdu,
       pch = 2) #plot character 2, triangles

```



Here we only have 1 call to a new plot object. The functions `lines()` and `points()` will place their output on top of the current plot (as long as they are in scale!). Some data types even have their own plot processes that run when you try to plot them. We will need a couple more packages, though, to deal with spatial data.

```
install.packages(c("rgdal", "sp", "geojsonio", "ggplot2", "leaflet"))
```

```

library(rgdal)
library(sp)
library(geojsonio)

```

```
## Warning: package 'geojsonio' was built under R version 3.6.3
```

```
library(ggplot2)
library(leaflet)
```

The following series of commands will be doing a few things:

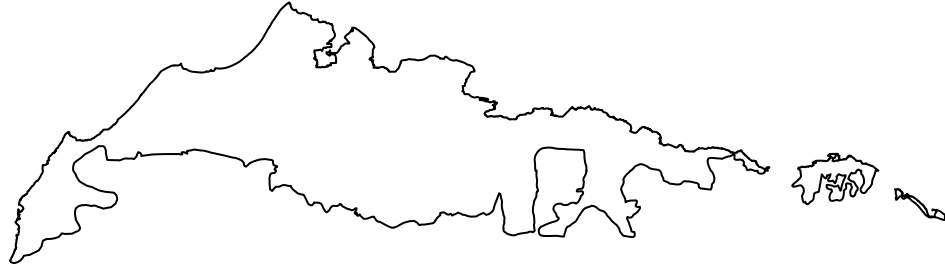
1. Download the PlotData.zip container from my GitHub page
2. Place the file in your current working directory
3. Unzip the files from the container
4. Delete the .zip file
5. Read in the ACP ESRI shp file
6. Project the file in WGS84
7. Plot the polygons using base R plot()

```
download.file("http://github.com/cfrost3/MBM_R_Short_Course/raw/master/PlotData/PlotData.zip",
              destfile = "PlotData.zip" , mode='wb')
unzip("PlotData.zip", exdir = "./Plot")
file.remove("PlotData.zip")
```

```
map= readOGR("./Plot/ACP_2007to2019_StudyArea.shp",
             layer = "ACP_2007to2019_StudyArea",
             verbose=FALSE)

map.proj <- spTransform(map, "+proj=longlat +ellps=WGS84 +datum=WGS84")

plot(map.proj)
```

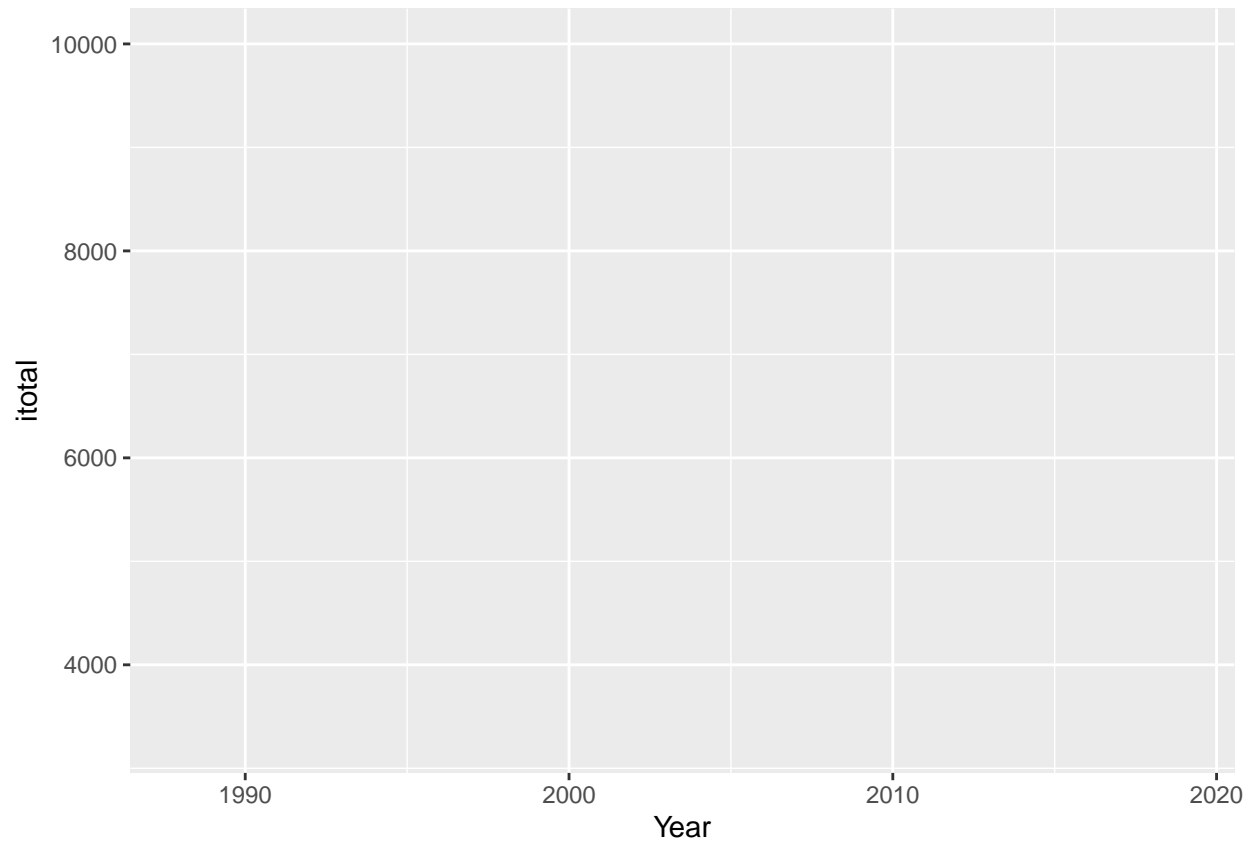


This is definitely the dry white toast of GIS mapping. We could make it work if we had to, but luckily we have plenty of added features in the contributed packages `ggplot2` and `leaflet` to visualize our spatial (and time series, and virtually every other possible type of) data.

4.2.2 Plotting in Contributed Packages

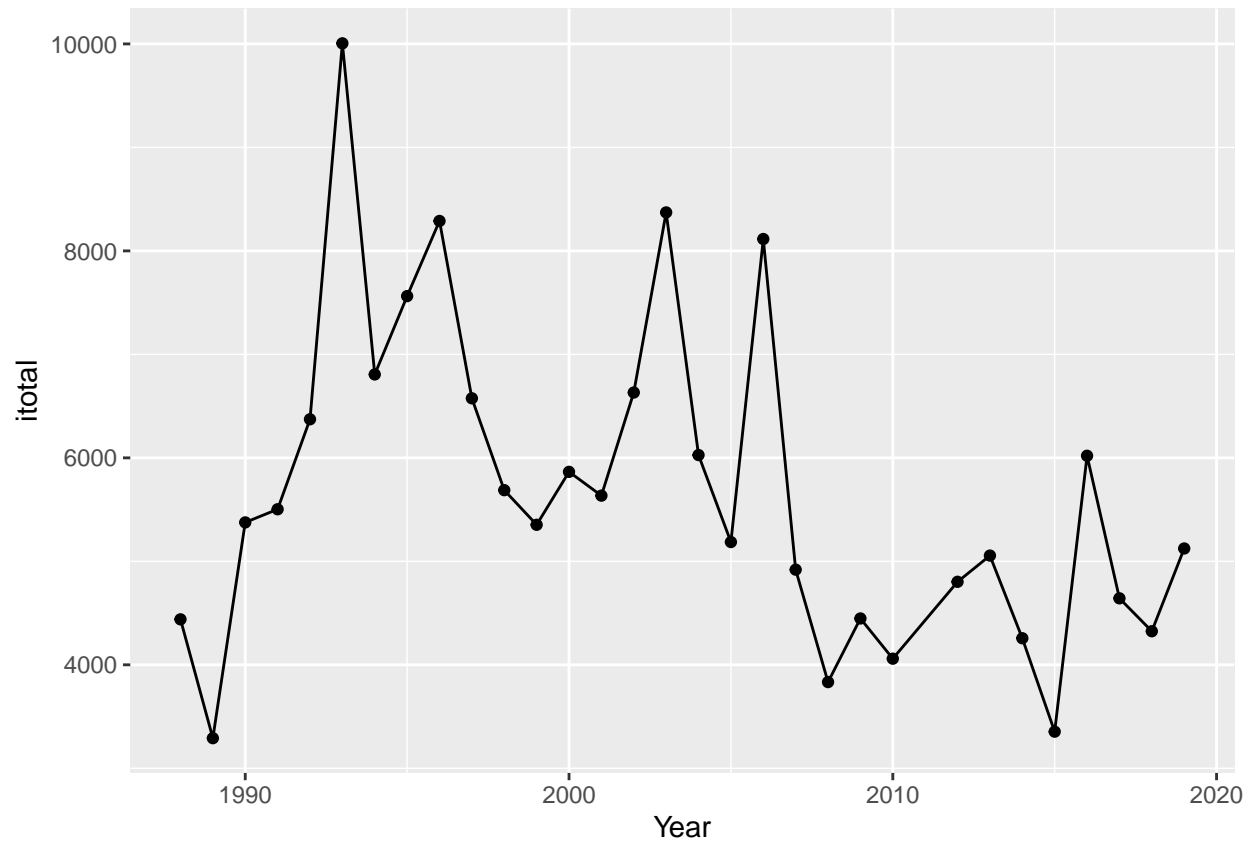
The best package (currently) to learn to plot in is `ggplot2`. It works with tidy data to produce publication-quality figures. The major difference in base R plotting and `ggplot2` is that while in base R you are passing vectors, in `ggplot2` you use whole data frames. Individual variables can be extracted from the data frame in various ways to add feature layers to your `ggplot`. We will start by recreating the long-tailed duck plot from earlier.

```
YKDHistoric$combined %>%  
  filter(Species == "LTDU") %>%  
  ggplot(aes( x = Year, y = itotal))
```



Mind-blowing! But really, what we've done here is establish our base layer and our aesthetics. Here, the x axis defined as Year, and the y axis as indicated total. We haven't told ggplot how to actual show us the data, but since the function is modular (like designing a modular home), we can start to tack on our layers.

```
YKDHistoric$combined %>%  
  filter(Species == "LTDU") %>%  
  ggplot(aes( x = Year, y = itotal)) +  
    geom_point() +  
    geom_line()
```

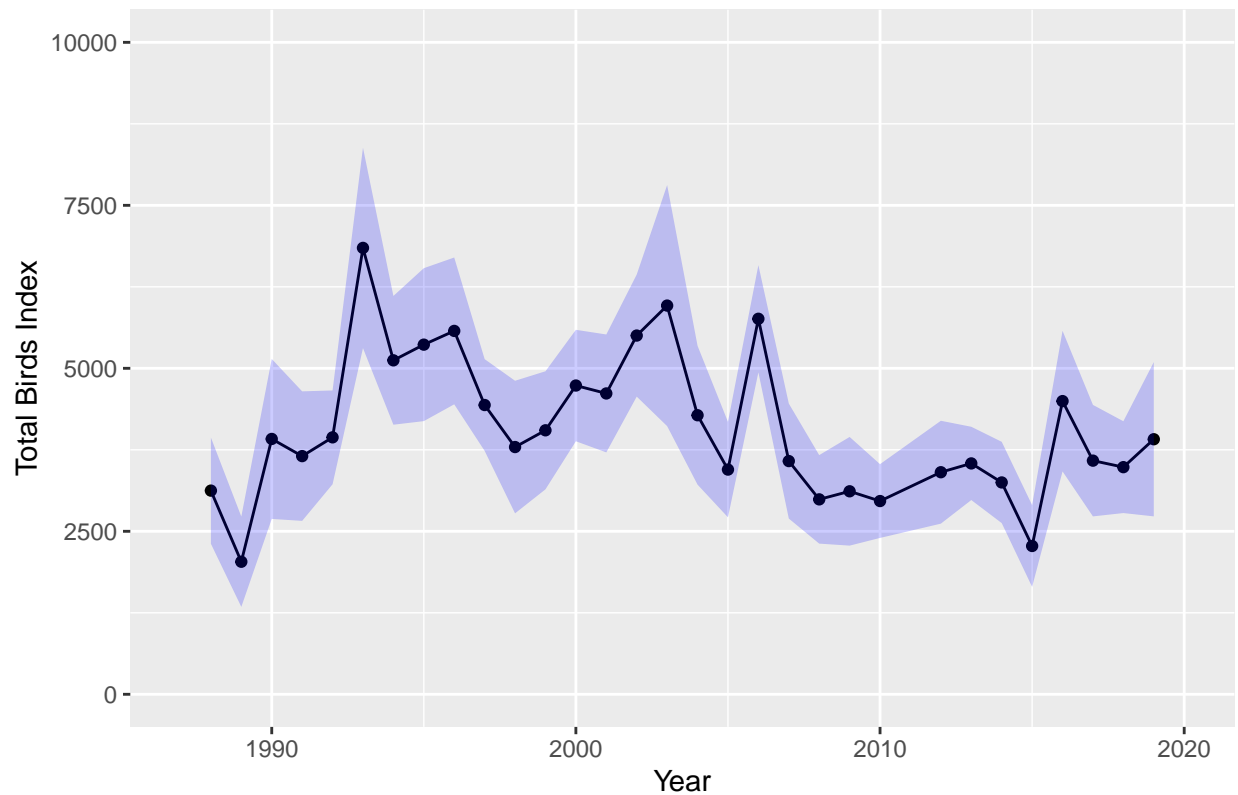


Here we told ggplot to now display our data as both points (`geom_point`) and lines (`geom_line`). Now we can fix the axes and get those confidence intervals on there. We will use a fill color and alpha value to set the transparency. We can also set the labels using `labs()` while we are at it (very intuitively).

```
duck.plot =
  YKDHistoric$combined %>%
  filter(Species == "LTDU") %>%
  ggplot(aes( x = Year, y = total)) +
    geom_point() +
    geom_line() +
    geom_ribbon(aes(ymin = total - 1.96 * total.se, ymax = total + 1.96 * total.se),
              fill = "blue",
              alpha = 0.2) +
    coord_cartesian(xlim=c(1987,2020), ylim=c(0, 10000)) +
    labs(title="LTDU Total Birds Index on the YK Delta, 1986-2019", x="Year", y="Total Birds Index")

print(duck.plot)
```


LTDU Total Birds Index on the YK Delta, 1986–2019



So far so good. Since a ggplot functions as an R object, we can store it in R (here we saved it as `duck.plot`) and add onto it with ease. The only difference is when we store the object it won't automatically plot, so we have to explicitly make it `print()`. Now let's imagine we want to throw a simple linear regression onto it.

```
duck.plot + geom_smooth(method='lm')
```

LTDU Total Birds Index on the YK Delta, 1986–2019



Pretty slick. Let's see that killer table we made earlier now in plot form.

```
emgo =  
  
  #tidy it up and pivot the estimates out  
  
  YKGHistoric$combined %>%  
  
  filter(Species=="EMGO") %>%  
  
  select(Year, itotal, itotal.se) %>%  
  
  mutate(avg3=rollapply(itotal,3,mean,align='right',fill=NA)) %>%  
  
  mutate(avg10=rollapply(itotal,10,mean,align='right',fill=NA)) %>%  
  
  pivot_longer(cols=c("itotal", "avg3", "avg10"), names_to = "index", values_to = "estimate")  
  
  #now plot it  
  
  emgo.base =  
  
  ggplot(emgo, aes(x = Year)) +  
  
  #change the linetype (dash spacing) by index type
```

```

geom_line(aes(y = estimate, linetype=index), size = 1) +

#add the confidence intervals
#note these are specific to itotal, so we have to filter what we pass

geom_ribbon(data = filter(emgo, index=="itotal"), aes(ymin = estimate - 1.96 * itotal.se, ymax = estimate + 1.96 * itotal.se),
            fill = "blue",
            alpha = 0.2) +

coord_cartesian(xlim=c(1985,2020), ylim=c(0, 40000)) +

labs(title="Indicated Total EMGO on the YK Delta, 1986-2019", x="Year", y="Indicated Total") +

#make a keen legend to sort out the mess

scale_linetype_manual(name = "Estimate",
                      values = c(1,2,3),
                      labels = c("Annual Estimate", "3-year Avg", "10-year Avg"),
                      limits = c("itotal", "avg3", "avg10")) +

#take the gray tiles off of the back

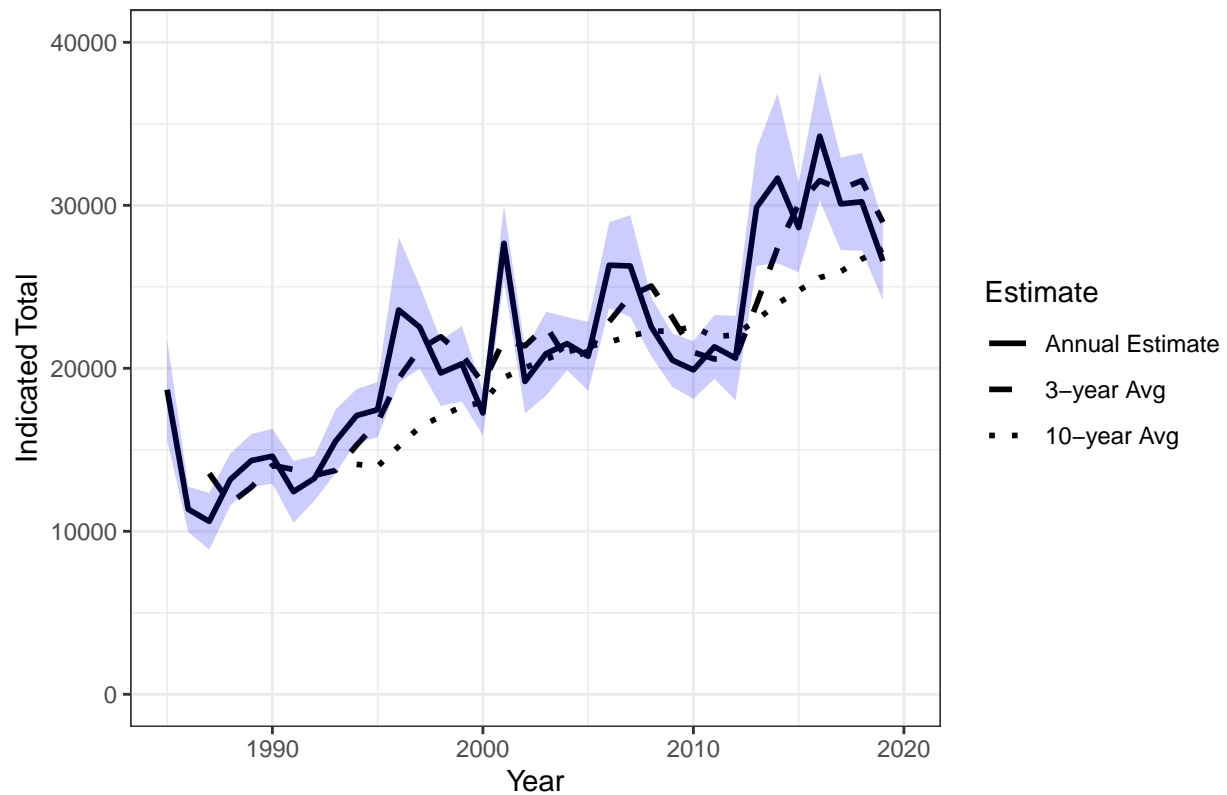
theme_bw()

print(emgo.base)

```

```
## Warning: Removed 11 row(s) containing missing values (geom_path).
```


Indicated Total EMGO on the YK Delta, 1986–2019



One powerful feature of ggplot is its ability to handle factors (or character strings that can be easily converted to factors). It can process a multi-species or multi-area data set and produce a string of plots in no time.

```
multi =

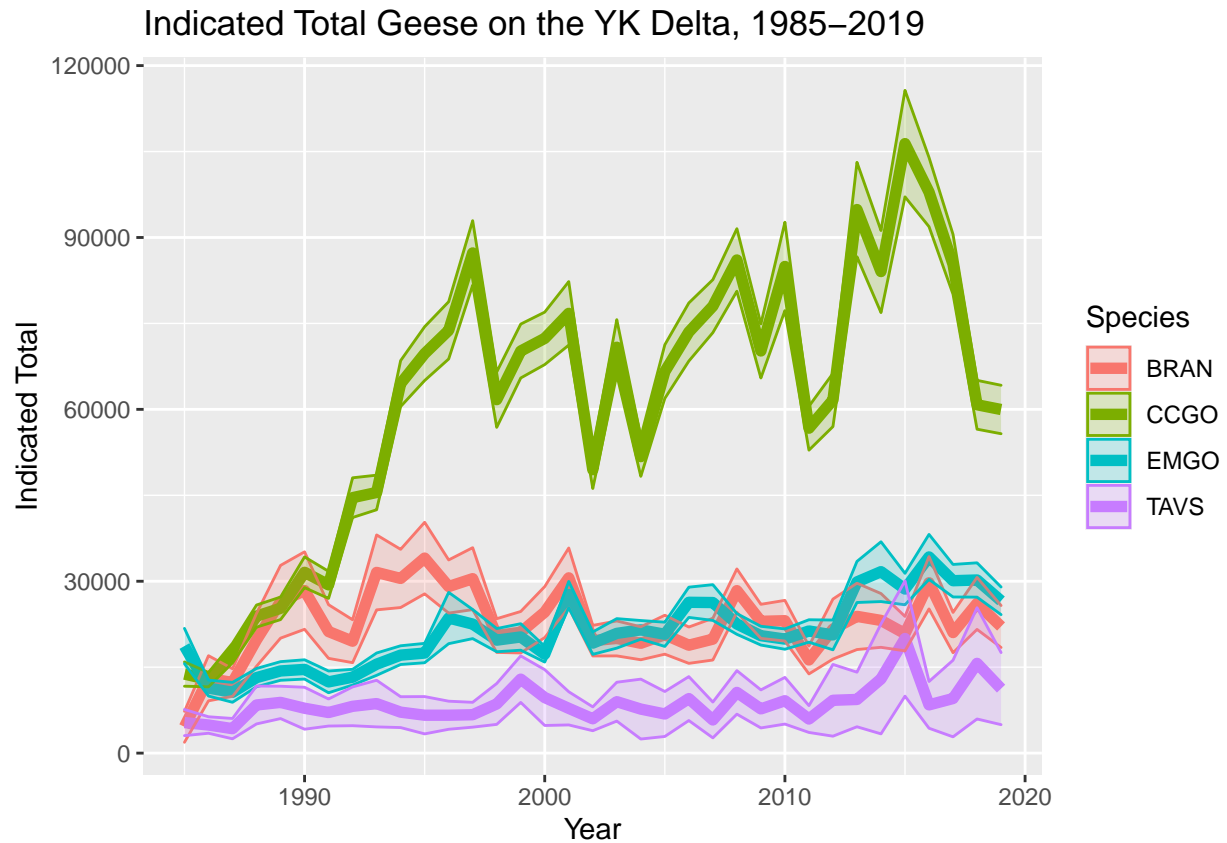
YKGHistoric$combined %>%

filter(Species %in% c("EMGO", "BRAN", "CCGO", "TAVS")) %>%

select(Year, Species, itotal, itotal.se)

multi.plot =
  ggplot(multi, aes(x = Year, y = itotal, color = Species)) +
  geom_line(size=2) +
  geom_ribbon(aes(ymin = itotal - 1.96 * itotal.se, ymax = itotal + 1.96 * itotal.se,
                  color=Species,
                  fill = Species),
            alpha=.2) +
  labs(title="Indicated Total Geese on the YK Delta, 1985-2019", x="Year", y="Indicated Total")

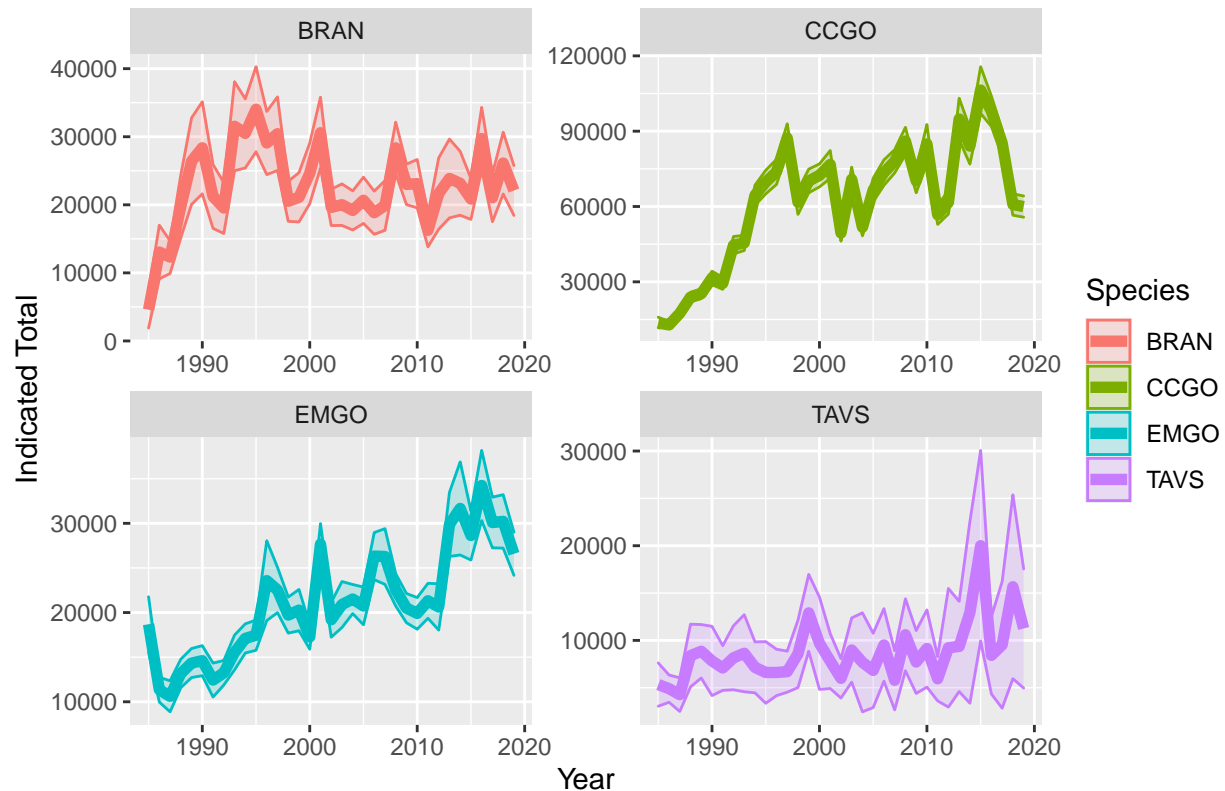
print(multi.plot)
```



The neat part here is that when we have a factor (or a character that can be automatically converted to one, in this case) such as `Species`, we can explicitly tell `ggplot` to map it as an aesthetic in the `aes()` definition right away. Then later when we want to change colors, linetypes, sizes, etc., we just tell it to change per our aesthetic. We even get a legend made automatically. But what if we wanted to split these out by our `Species` factor, each to their own figure?

```
multi.plot +  
  facet_wrap(~Species, scales = "free")
```

Indicated Total Geese on the YK Delta, 1985–2019



We told it to create separate figures (facets) wrapped logically on our screen based on how big or small we shape our plot window. These were separated by species (`~Species`), and we had to tell ggplot to allow the scales for each species to be “free” and not all based on the highest value of the most abundant species. The alternative (and default) is “fixed.” Below we will use `facet_grid` to tell ggplot to change the number of rows to reflect the number of Species it counts. I also use a few more style elements here to expand the x axis values and produce the legend.

```
ssm = read.csv("./Plot/StateSpaceEst.csv",
               header = TRUE,
               stringsAsFactors = FALSE)

ssm = ssm %>%
  pivot_longer(cols=c("N", "index"),
               names_to = "method",
               values_to = "estimate")

ssm.plot=
  ggplot(ssm, aes(x = Year, y = estimate, fill = method)) +
  geom_line(data = ssm %>% filter(method == "index"), size=1) +
  geom_point(data = ssm %>% filter(method == "index"), size=2) +
  geom_ribbon(data = ssm %>% filter(method == "index"),
             aes(ymin = l95index, ymax = u95index),
             alpha = 0.6) +

  geom_ribbon(data = ssm %>% filter(method == "N"),
             aes(ymin = lower95, ymax = upper95),
```

```

alpha = 0.3) +

facet_grid(Species~., scales = "free") +

labs(title="Indicated Total on the ACP, 2009–2018,\nIncluding State-Space Model Estimates",
      x="Year", y="Index / Model Estimate") +

scale_fill_manual(name = "Method\n(95% CI)",
                  values=c("grey", "red") ,
                  labels = c("Index Estimate", "State Space Estimate")) +

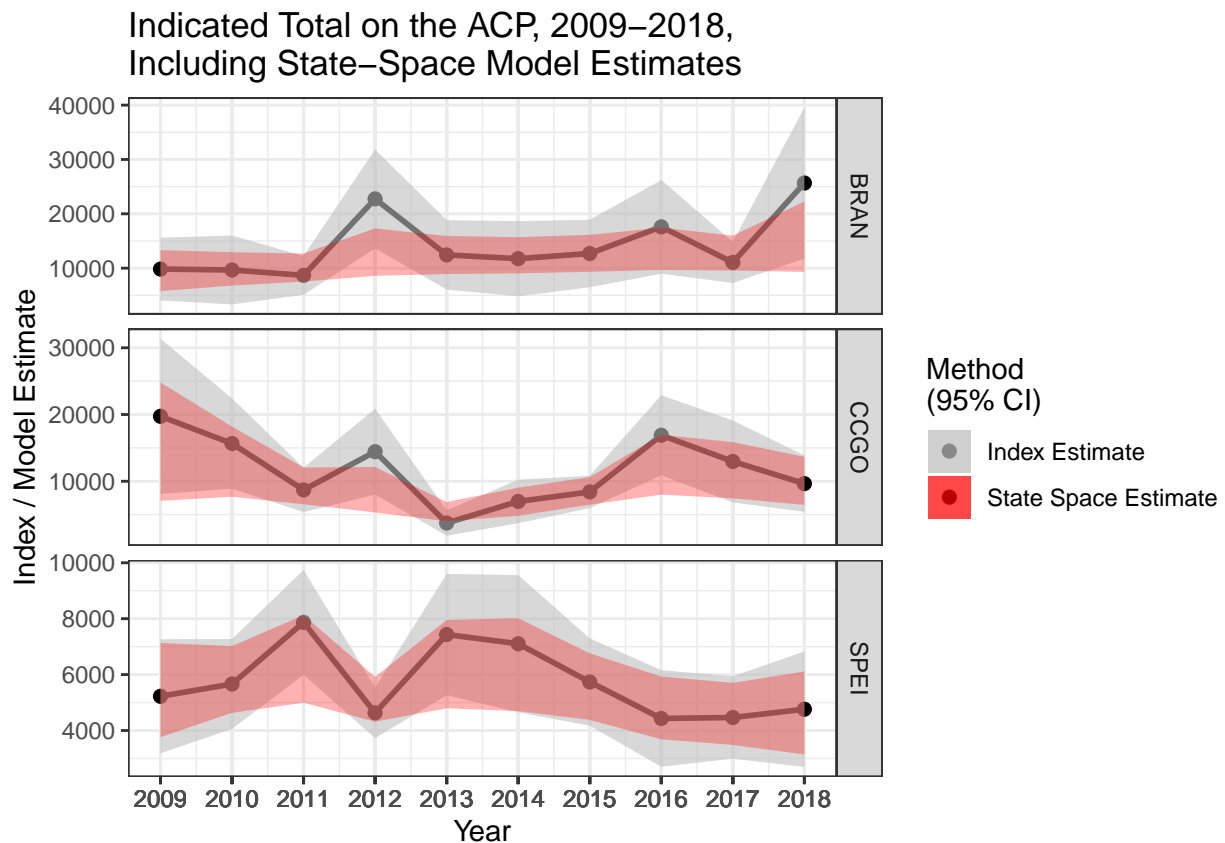
scale_x_continuous("Year", labels = as.character(ssm$Year), breaks = ssm$Year) +

theme(axis.text.x = element_text(face="bold", color="black")) +

theme_bw()

print(ssm.plot)

```



Or, for a single species:

```

ssm.single = ssm %>% filter(Species == "BRAN")

ssm.single.plot =

```

```

ggplot(data = ssm.single,
       aes(x = Year, y = estimate, fill = method)) +

geom_line(data = ssm.single %>% filter(method == "index"), size=1) +

geom_point(data = ssm.single %>% filter(method == "index"), size=2) +

geom_ribbon(data = ssm.single %>% filter(method == "index"),
          aes(ymin = l95index, ymax = u95index),
          alpha = 0.6) +

geom_ribbon(data = ssm.single %>% filter(method == "N"),
          aes(ymin = lower95, ymax = upper95),
          alpha = 0.3) +

labs(title="Indicated Total Brant on the ACP, 2009-2018,\nIncluding State-Space Model Estimates",
     x="Year", y="Index / Model Estimate") +

scale_fill_manual(name = "Method\n(95% CI)",
                  values=c("grey", "red") ,
                  labels = c("Index Estimate", "State Space Estimate")) +

scale_x_continuous("Year",
                  labels = as.character(ssm.single$Year),
                  breaks = ssm.single$Year) +

theme(axis.text.x = element_text(face="bold", color="black")) +

coord_cartesian(ylim=c(0, 1.1 * max(ssm.single$u95index))) +

geom_text(x=min(ssm.single$Year + 1), y=max(ssm.single$u95index),
          label=paste0("Mean r = ", round(1 + ssm.single$mean.r[1], 2))) +

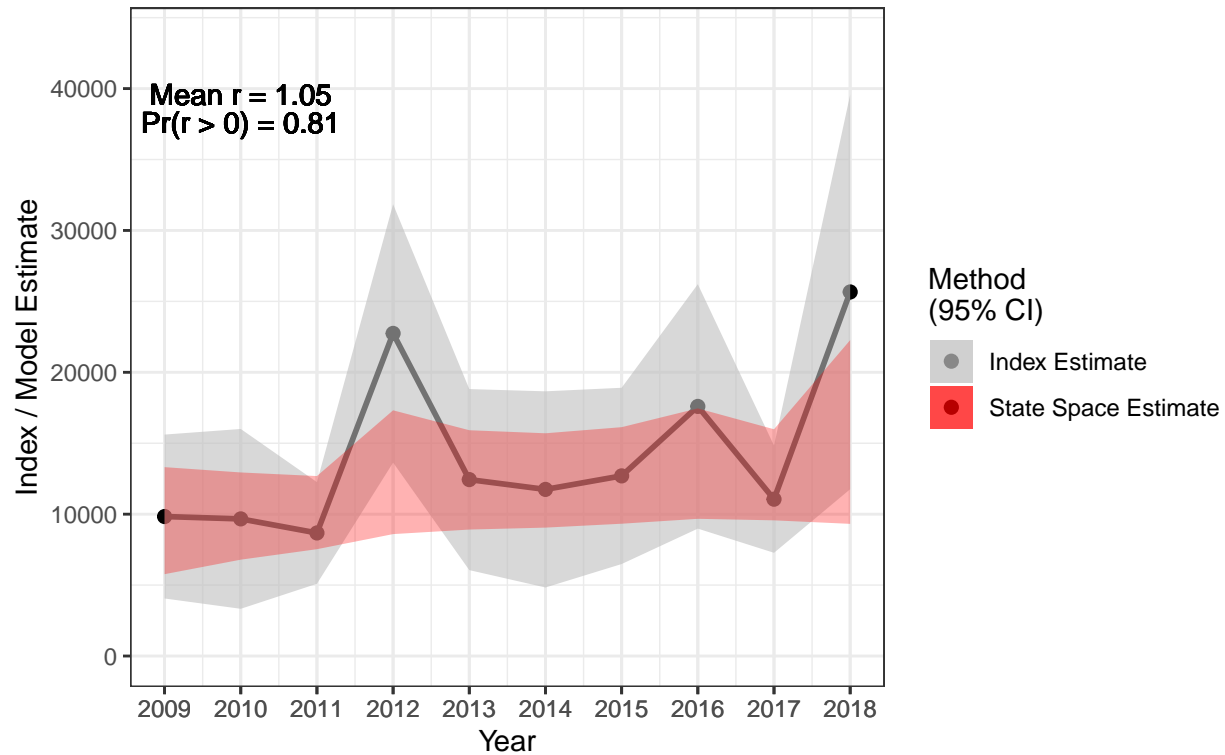
geom_text(x=min(ssm.single$Year + 1), y=.95*max(ssm.single$u95index),
          label=paste0("Pr(r > 0) = ", round(ssm.single$p.r, 2))) +

theme_bw()

print(ssm.single.plot)

```

Indicated Total Brant on the ACP, 2009–2018, Including State–Space Model Estimates



Note the use of `geom_text` to paste together a character string and place it at a given (x,y) on the figure. Also take note of the way you can dynamically set an axis to be a little wider or more narrow than the automatic setting. We can also quickly add the associated table.

Table 3: Indicated Total Brant on the ACP, 2009–2018, including state-space model estimates.

Year	Indicated Total*	95% CI		State-Space Estimate	95% CI	
		Lower	Upper		Lower	Upper
2009	9833	4057	15610	9393	5769	13315
2010	9667	3329	16005	9724	6792	12941
2011	8676	5095	12256	10118	7531	12687
2012	22743	13638	31848	11826	8595	17321
2013	12437	6056	18817	11869	8915	15916
2014	11746	4831	18660	12048	9051	15702
2015	12698	6487	18910	12460	9327	16133
2016	17597	8983	26210	13088	9673	17456
2017	11055	7278	14832	12766	9566	15987
2018	25663	11760	39566	14567	9316	22271

* Indicated Total is calculated as 2(singles + pairs) + flocked birds.

Can ggplot plot spatial information? It can, but not necessarily easily. Remember when I said that ggplot takes a data frame rather than a vector? Well, it also takes a data frame rather than a spatial format. So to make it happen, we have to convert a shapefile or other spatial overlay into a data frame. It isn't too hairy and I've created a function (`LoadMap`) in `AKAerial` that will read in your spatial file and output it as a ggplot-able data frame. We will use it here and send it our shapefile.

```
map.df = AKAerial::LoadMap("./Plot/ACP_2007to2019_StudyArea.shp")


ggplot() +

  #geom_path here tells ggplot to connect the dots of my spatial file,
  #since polygons in their simplest form are just a series of vertices.

  geom_path(data=map.df, aes(long,lat,group=group)) +

  geom_path(color="black") +

  #now set a reasonable bounding box
  coord_map(xlim=c(min(map.df$long), max(map.df$long)),
            ylim=c(min(map.df$lat), max(map.df$lat)))
```




This isn't spectacular either, maybe a step up from base R spatial plots. We can add functionality to it and trace other polygons and fill in solid colors, but for spatial visualization, nothing really compares to leaflet (yet).

Leaflet is a powerful JavaScript library that creates user-friendly maps for almost any occasion. It isn't explicitly R code, so the package leaflet was created to access the JavaScript functionality through R. You can see how powerful it is with even a basic map.

```
leaflet() %>%
  addTiles() %>% # use the default base map tiles
  addMarkers(lng=-87.6553, lat=41.9484,
```

```
popup="Wrigley Field")
```

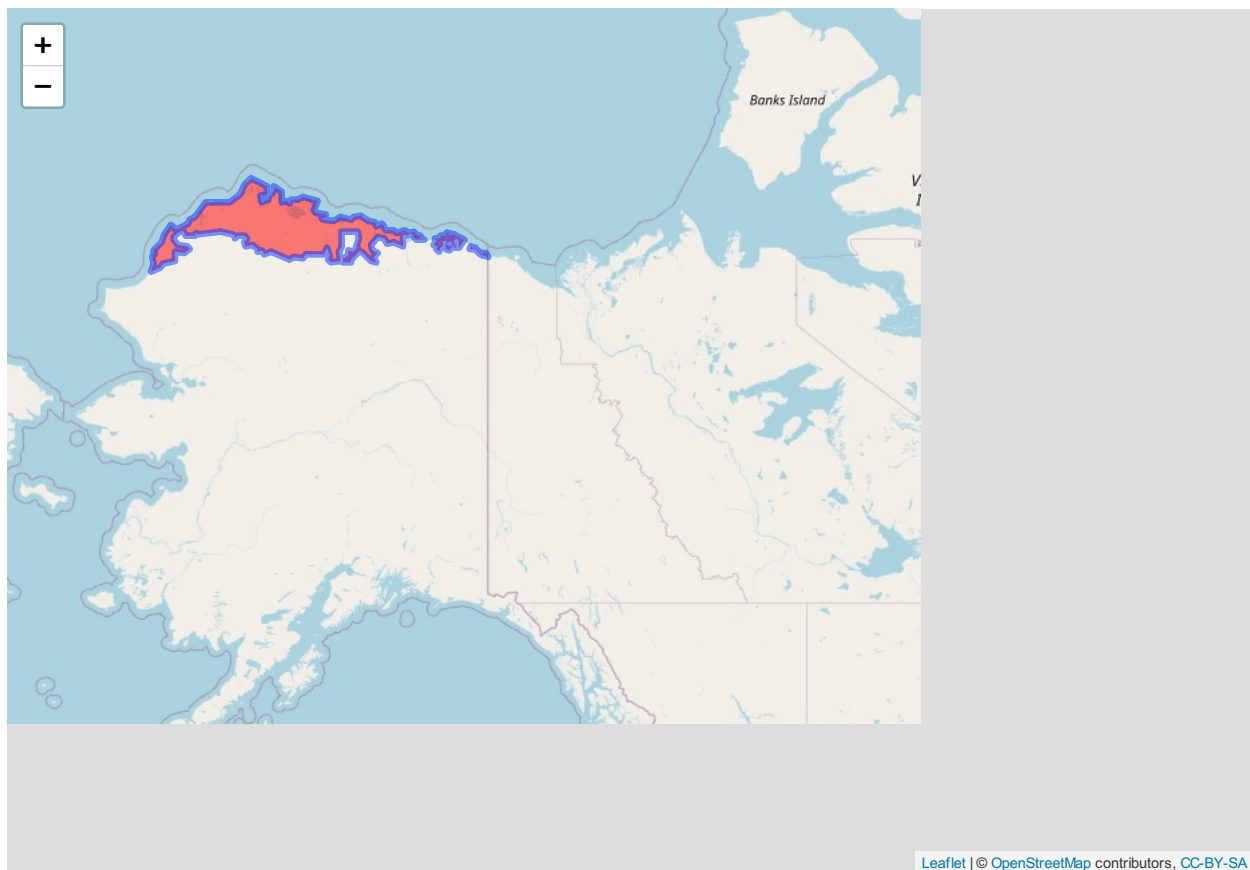


Leaflet | © OpenStreetMap contributors, CC-BY-SA

Since leaflet maps are html widgets, they won't appear well or function in this pdf document. On your computer screen, though, you can click, drag, and zoom. They also work great in interactive html documents (much like datatables and kables).

In addition to just displaying a place on a map, leaflet can take your spatial data files and make glorious visualizations with them.

```
leaflet() %>%  
  addTiles() %>%  
  addPolygons(data=map.proj,  
    fillColor = "red",  
    fillOpacity = .5)
```

We can add almost any spatial files to it.

```
#read and project our lines shp file

lines= readOGR("./Plot/ACP_2019_Transects.shp",
  layer = "ACP_2019_Transects",
  verbose=FALSE)

lines.proj <- spTransform(lines, "+proj=longlat +ellps=WGS84 +datum=WGS84")

#map it

leaflet() %>%

  addPolygons(data=map.proj,
    color = "yellow",
    fill = FALSE,
    fillOpacity = .5) %>%

  #add our lines here
  addPolylines(data=lines.proj,
    color="white",
    weight=4,
    opacity=.9,
    label=~TRANSID,
    popup = paste("Transect: ", lines.proj$TRANSID, "<br>",
```

```

"Length: ", lines.proj$LENGTH)) %>%

#scale for...scale
addScaleBar() %>%

#cool satellite imagery base map
addProviderTiles("Esri.WorldImagery") %>%

#add labels
addLabelOnlyMarkers(data = fortify(lines.proj) %>%
  filter(order == 1) %>%
  mutate(new.id = as.numeric(id) + 1),
  label = ~as.character(new.id),
  labelOptions = labelOptions(noHide = T,
                                direction = 'top',
                                textOnly = T,
                                style = list(
                                  "color" = "white",
                                  "font-size" = "12px"))))

```



Leaflet | Tiles © Esri — Source: Esri, i-cubed, USDA, USGS, AEX, GeoEye, Getmapping, Aerogrid, IGN, IGP, UPR-EGP and the GIS User Community

We got a little fancy there by extracting the start point of each of our transect lines, then adding a text label to the point with the transect id. Now let's add some observation data:

```

#read our observations

obs = read.csv("./Plot/ACP_2019_QCObs_HWilson.csv",

```

```

        header=TRUE,
        stringsAsFactors = FALSE)

#define the spatial x,y for plotting
coordinates(obs)=~Lon+Lat

#map it

leaflet() %>%

  addPolygons(data=map.proj,
              color = "yellow",
              fill = FALSE,
              fillOpacity = .5) %>%

  #add our lines here
  addPolylines(data=lines.proj,
              color="white",
              weight=4,
              opacity=.9,
              label=~TRANSID,
              popup = paste("Transect: ", lines.proj$TRANSID, "<br>",
                           "Length: ", lines.proj$LENGTH)) %>%

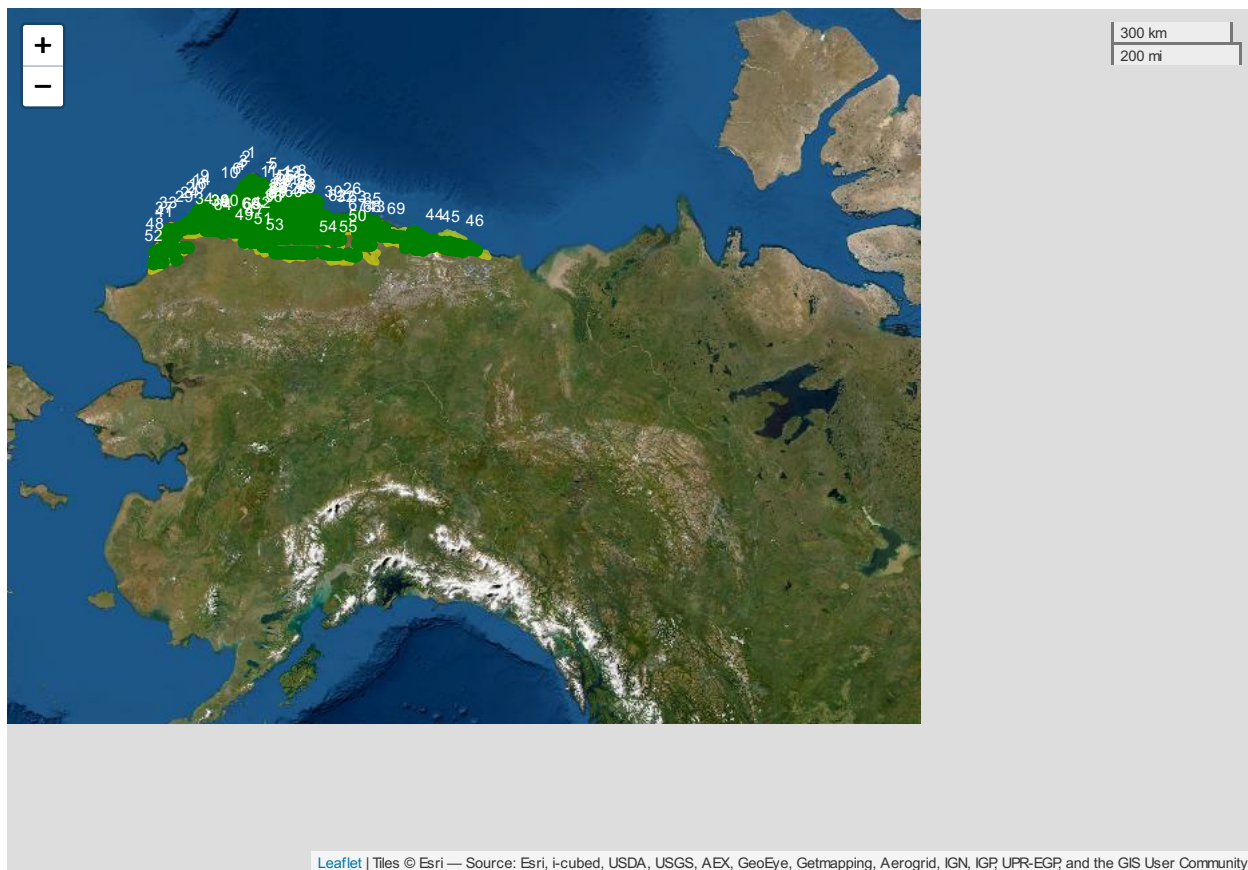
  #scale for...scale
  addScaleBar() %>%

  #cool satellite imagery base map
  addProviderTiles("Esri.WorldImagery") %>%

  #add labels
  addLabelOnlyMarkers(data = fortify(lines.proj) %>%
                      filter(order == 1) %>%
                      mutate(new.id = as.numeric(id) + 1),
                      label = ~as.character(new.id),
                      labelOptions = labelOptions(noHide = T,
                                                    direction = 'top',
                                                    textOnly = T,
                                                    style = list(
                                                      "color" = "white",
                                                      "font-size" = "12px")))) %>%

  #add the point data
  addCircleMarkers(data=obs,
                  radius = 5,
                  color = "green",
                  stroke = FALSE,
                  fillOpacity = 1,
                  popup= paste(obs$Observer, "<br>", obs$Species,
                              "<br>", obs$Obs_Type, "<br>", "n = ", obs$Num,
                              "<br>", "Transect ", obs$Transect, "<br>"))

```



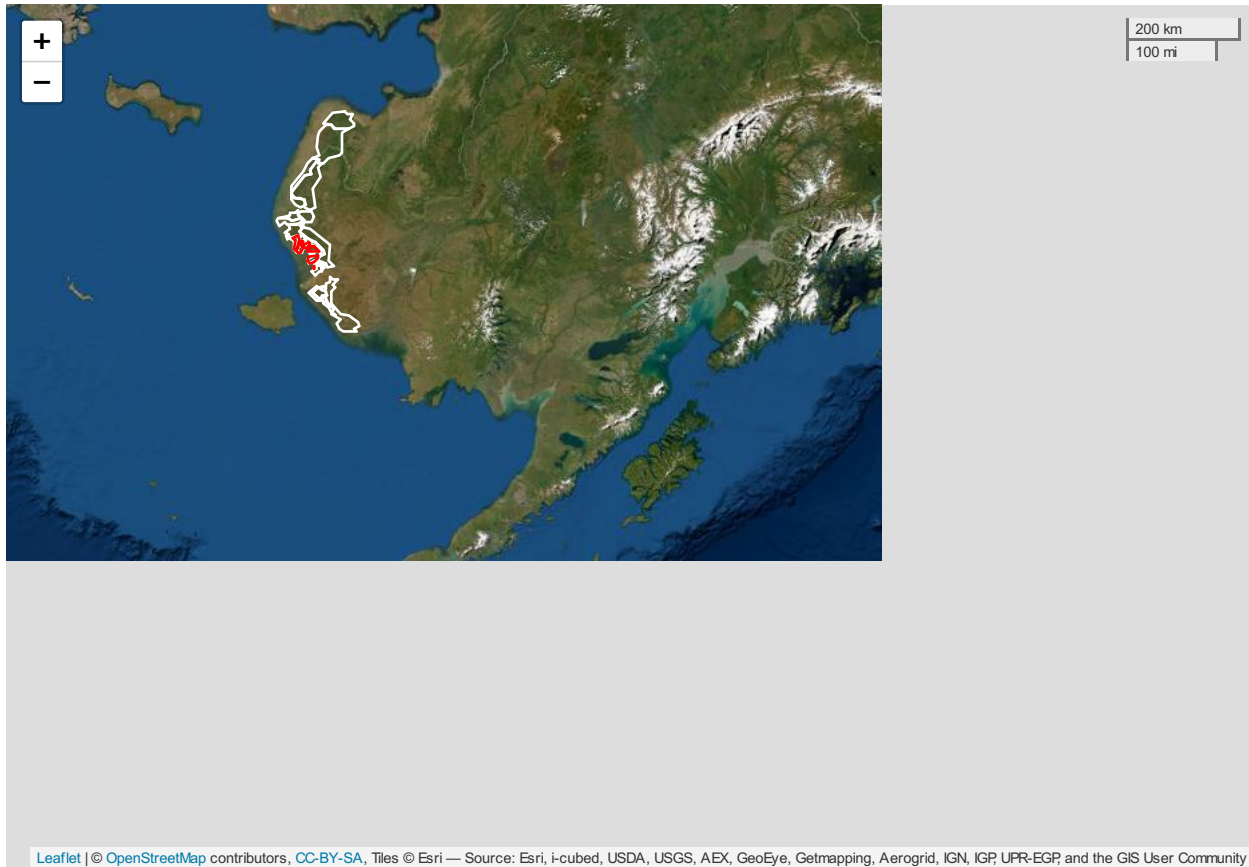
You can probably imagine many ways to use this type of map... at the very least, in QA/QC. For example, scroll to the east and look at all of the green points that fall out of the study area! In this case, though, that was actually by design for a corollary project on the refuge, but you can learn a lot from even a basic map. We can also accomplish all of this with GeoJSON files.

GeoJSON stands for Geographic Javascript Object Notation and is an encoding method for a variety of spatial data types. Here we will use some polygon GeoJSON containers from the Yukon-Kuskokwim Delta.

```
ykdair <- geojson_read("./Plot/YKD_DesignStrata.geojson", what = "sp")
ykdnest <- geojson_read("./Plot/NestPlotStudyAreaBoundary.geojson", what = "sp")

leaflet() %>%
  addTiles() %>%
  addPolygons(data=ykdair,
    fillColor="blue",
    fillOpacity=0,
    stroke=TRUE,
    color="white",
    opacity=1,
    weight=2) %>%
  addPolygons(data=ykdnest,
    fillOpacity=0,
    stroke=TRUE,
    color="red",
    opacity=1,
    weight=2) %>%
  addScaleBar() %>%
```

```
addProviderTiles("Esri.WorldImagery")
```



We read in 2 polygon layers describing the overall aerial survey coverage and subsetting nest plot project coverage. Zoom way in where the red and white overlap. Notice any potential QA/QC issues?